

# OS and Virtual Memory

CS110 Discussion 13

Suting Chen



# Topics for today

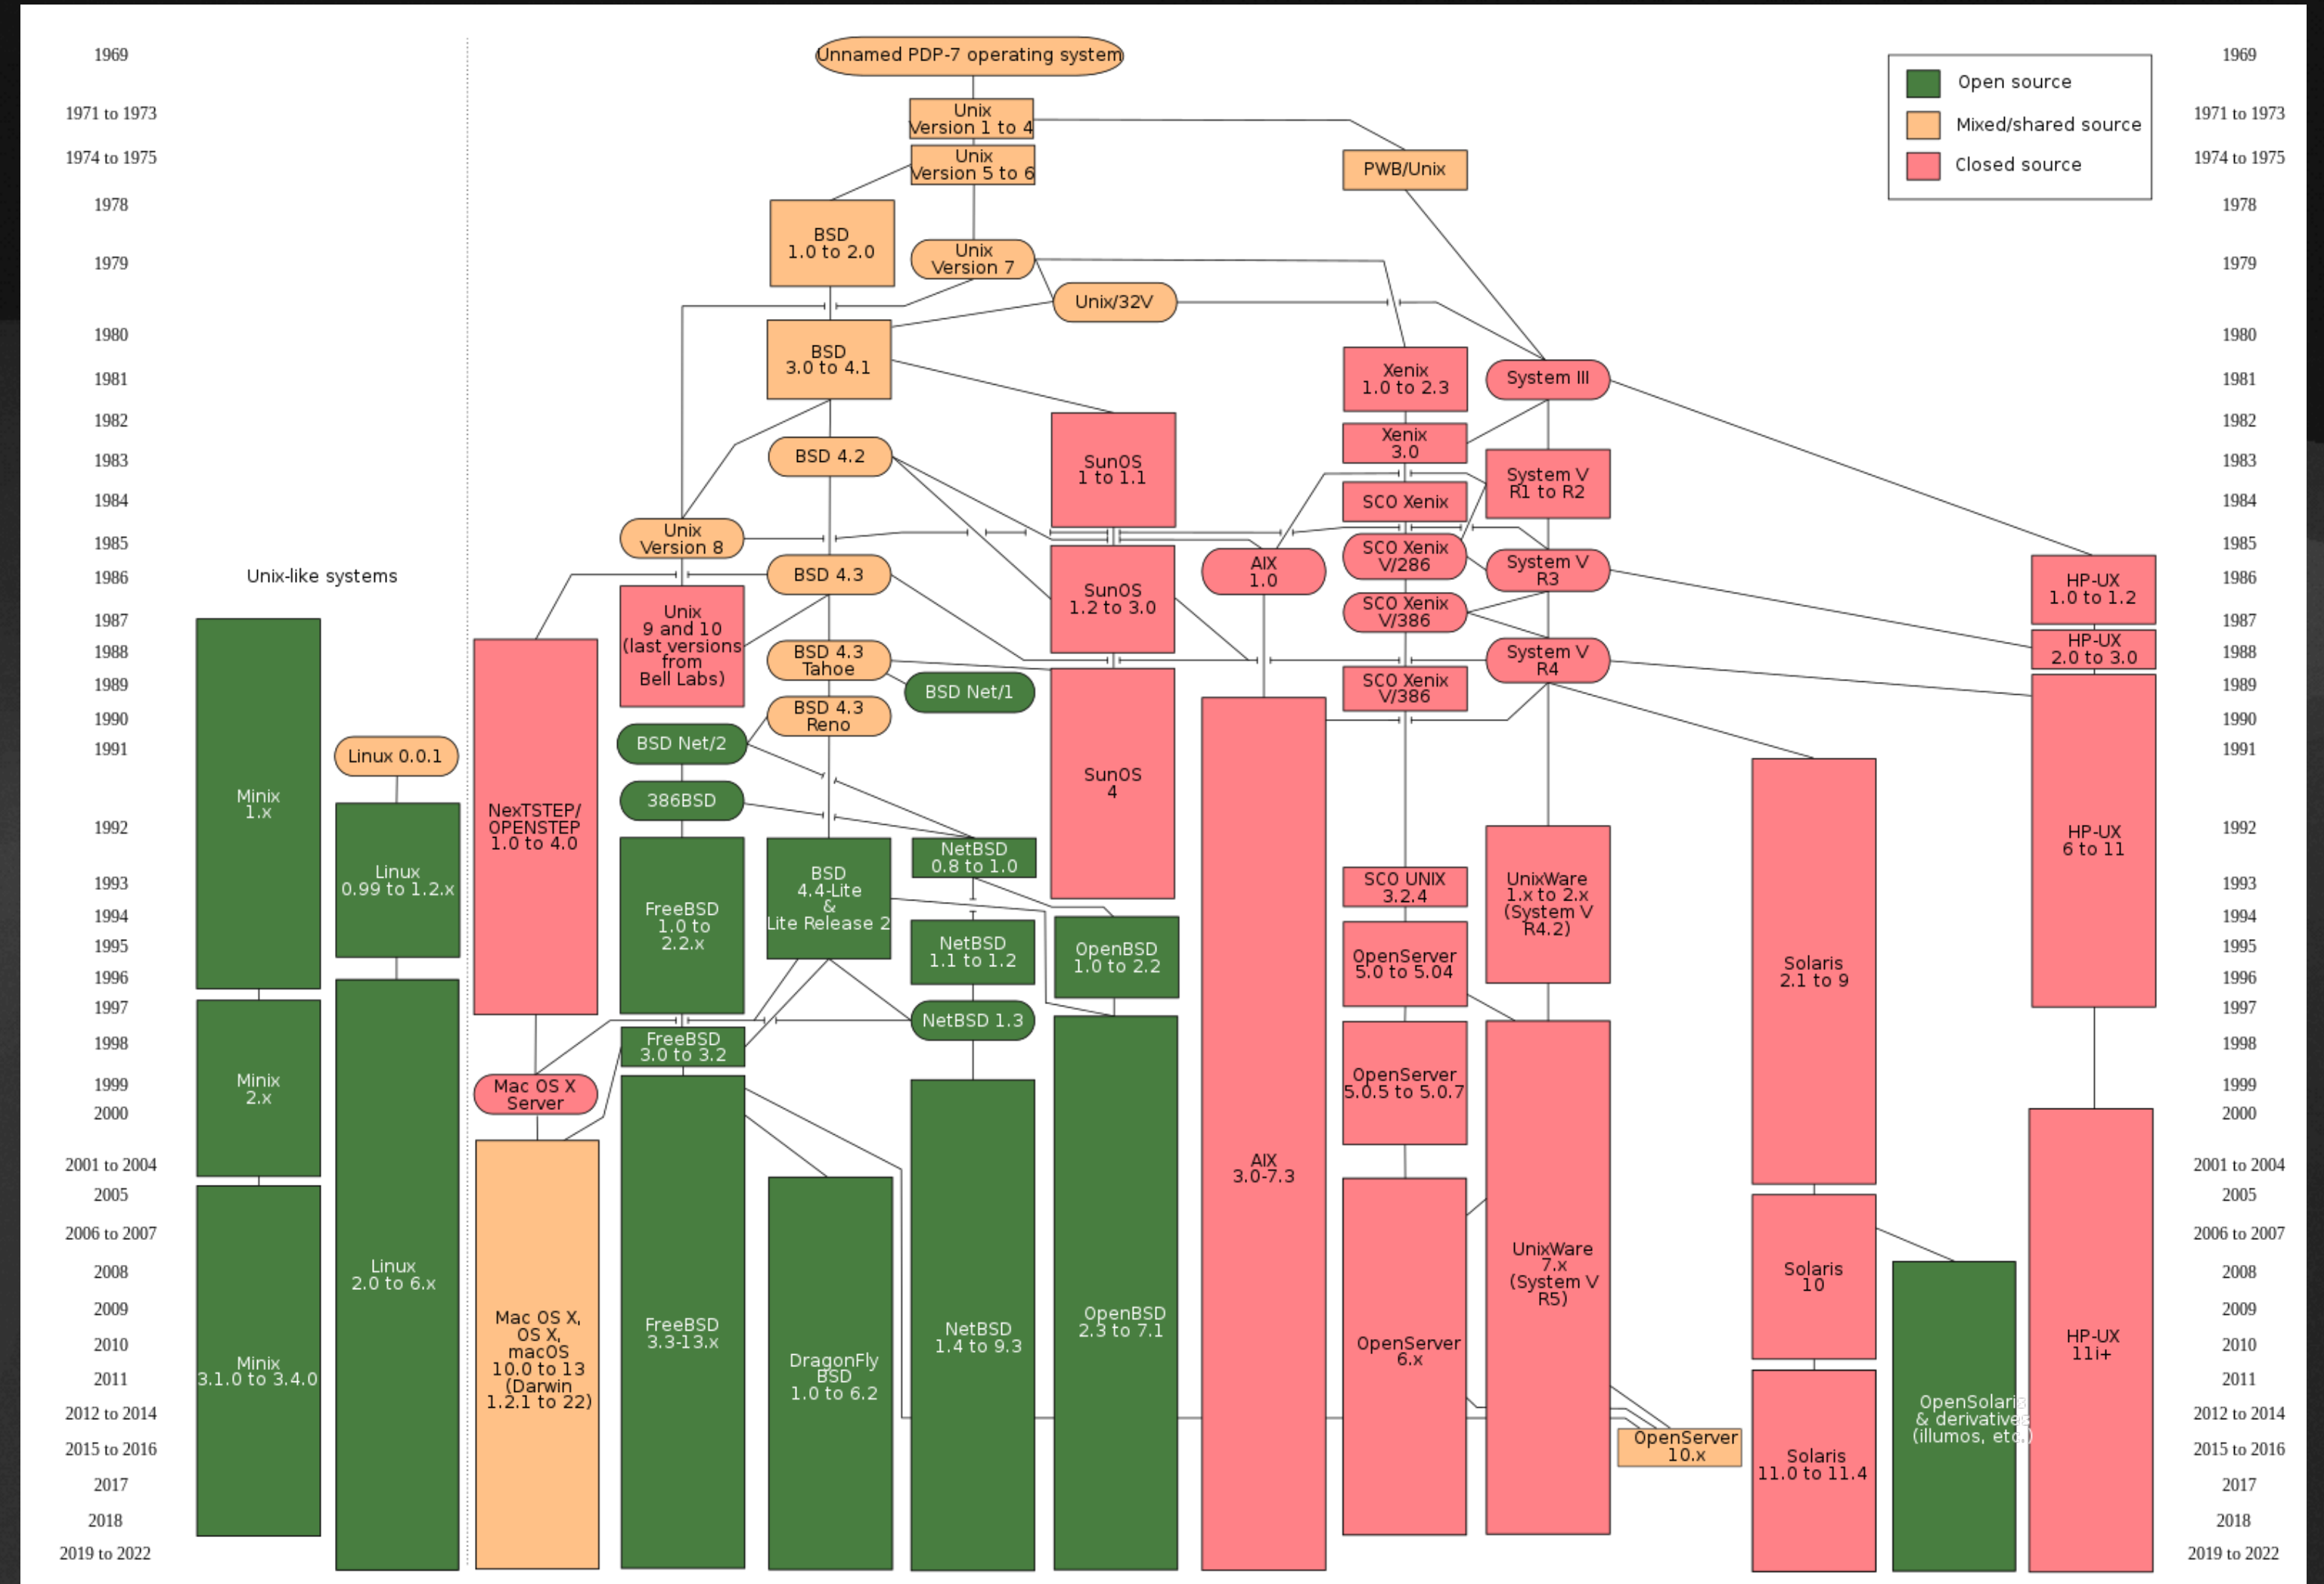
- What's Operating System?
- How OS helps your programs?
- Virtual Memory



# What's Operating System?

- Linux / Unix / Windows
- Debian / Fedora / Arch

# Suting Chen





# What Operating Systems do?

- Referee
- Glue
- Illusionist



# What Operating Systems do?

## Referee

- Protect your program from others
- Arrange system resources fairly



# What Operating Systems do?

## Glue

- Basic services: Loader (CALL)
- Provide hardware interfaces:
  - printf / scanf — Terminal
  - malloc — Memory
  - fopen — Disk
  - socket — Network card



# What Operating Systems do?

## Glue — Importance

- Saves programmers' time to read docs and interact with hardware
- Offers unified interface for multiple devices:
  - `pthread_create()` works for all platforms — Intel, AMD, Apple...
  - `fopen()` opens file on all devices — SSD, HDD, even floppy disks!
- Enables permissions

```
chmod 750 /home/cs110
```



# What Operating Systems do?


## Illusionist

- Each process believe that it have the full system resource, including:
  - Whole physical CPU core
  - Whole memory space
  - Direct access to any other resource as if there are no other processes



# Illusionist — Whole physical CPU core


When running single program...





# Illusionist — Whole physical CPU core

When running multiple programs ...





# Illusionist — Whole physical CPU core


When running multiple programs ...





# Illusionist — Whole physical CPU core

When running multiple programs ...





# Illusionist — Whole memory space


- Every program assert it has memory size of `0xFFFFFFFF`
- Problem:
  - We don't have so much memory to **waste**
  - We only have one `0x0` and one `0xFFFFFFFF`

`0x00000000`

code

static data

heap



`0xFFFFFFFF`

stack



# Illusionist — Whole memory space

If you have a 64-bit machine (certainly you do!)

- 64-bit addresses!
- Total size  $2^{64}$  Bytes = 16 EB =  $16 \times 1024$  PB
  - Impossible to find such memory (at least in 2023)

Suting Chen

实际上市面上的64位机器一般不会真的有64位的内存空间，这里为了演示方便忽略了很多问题  
具体可以去看AMD64 Architecture Programmer's Manual，你也可以用lscpu查看自己的CPU情况  
大部分情况下都是48位的Virtual Address Space (这里当故事听就好)



# Virtual Memory

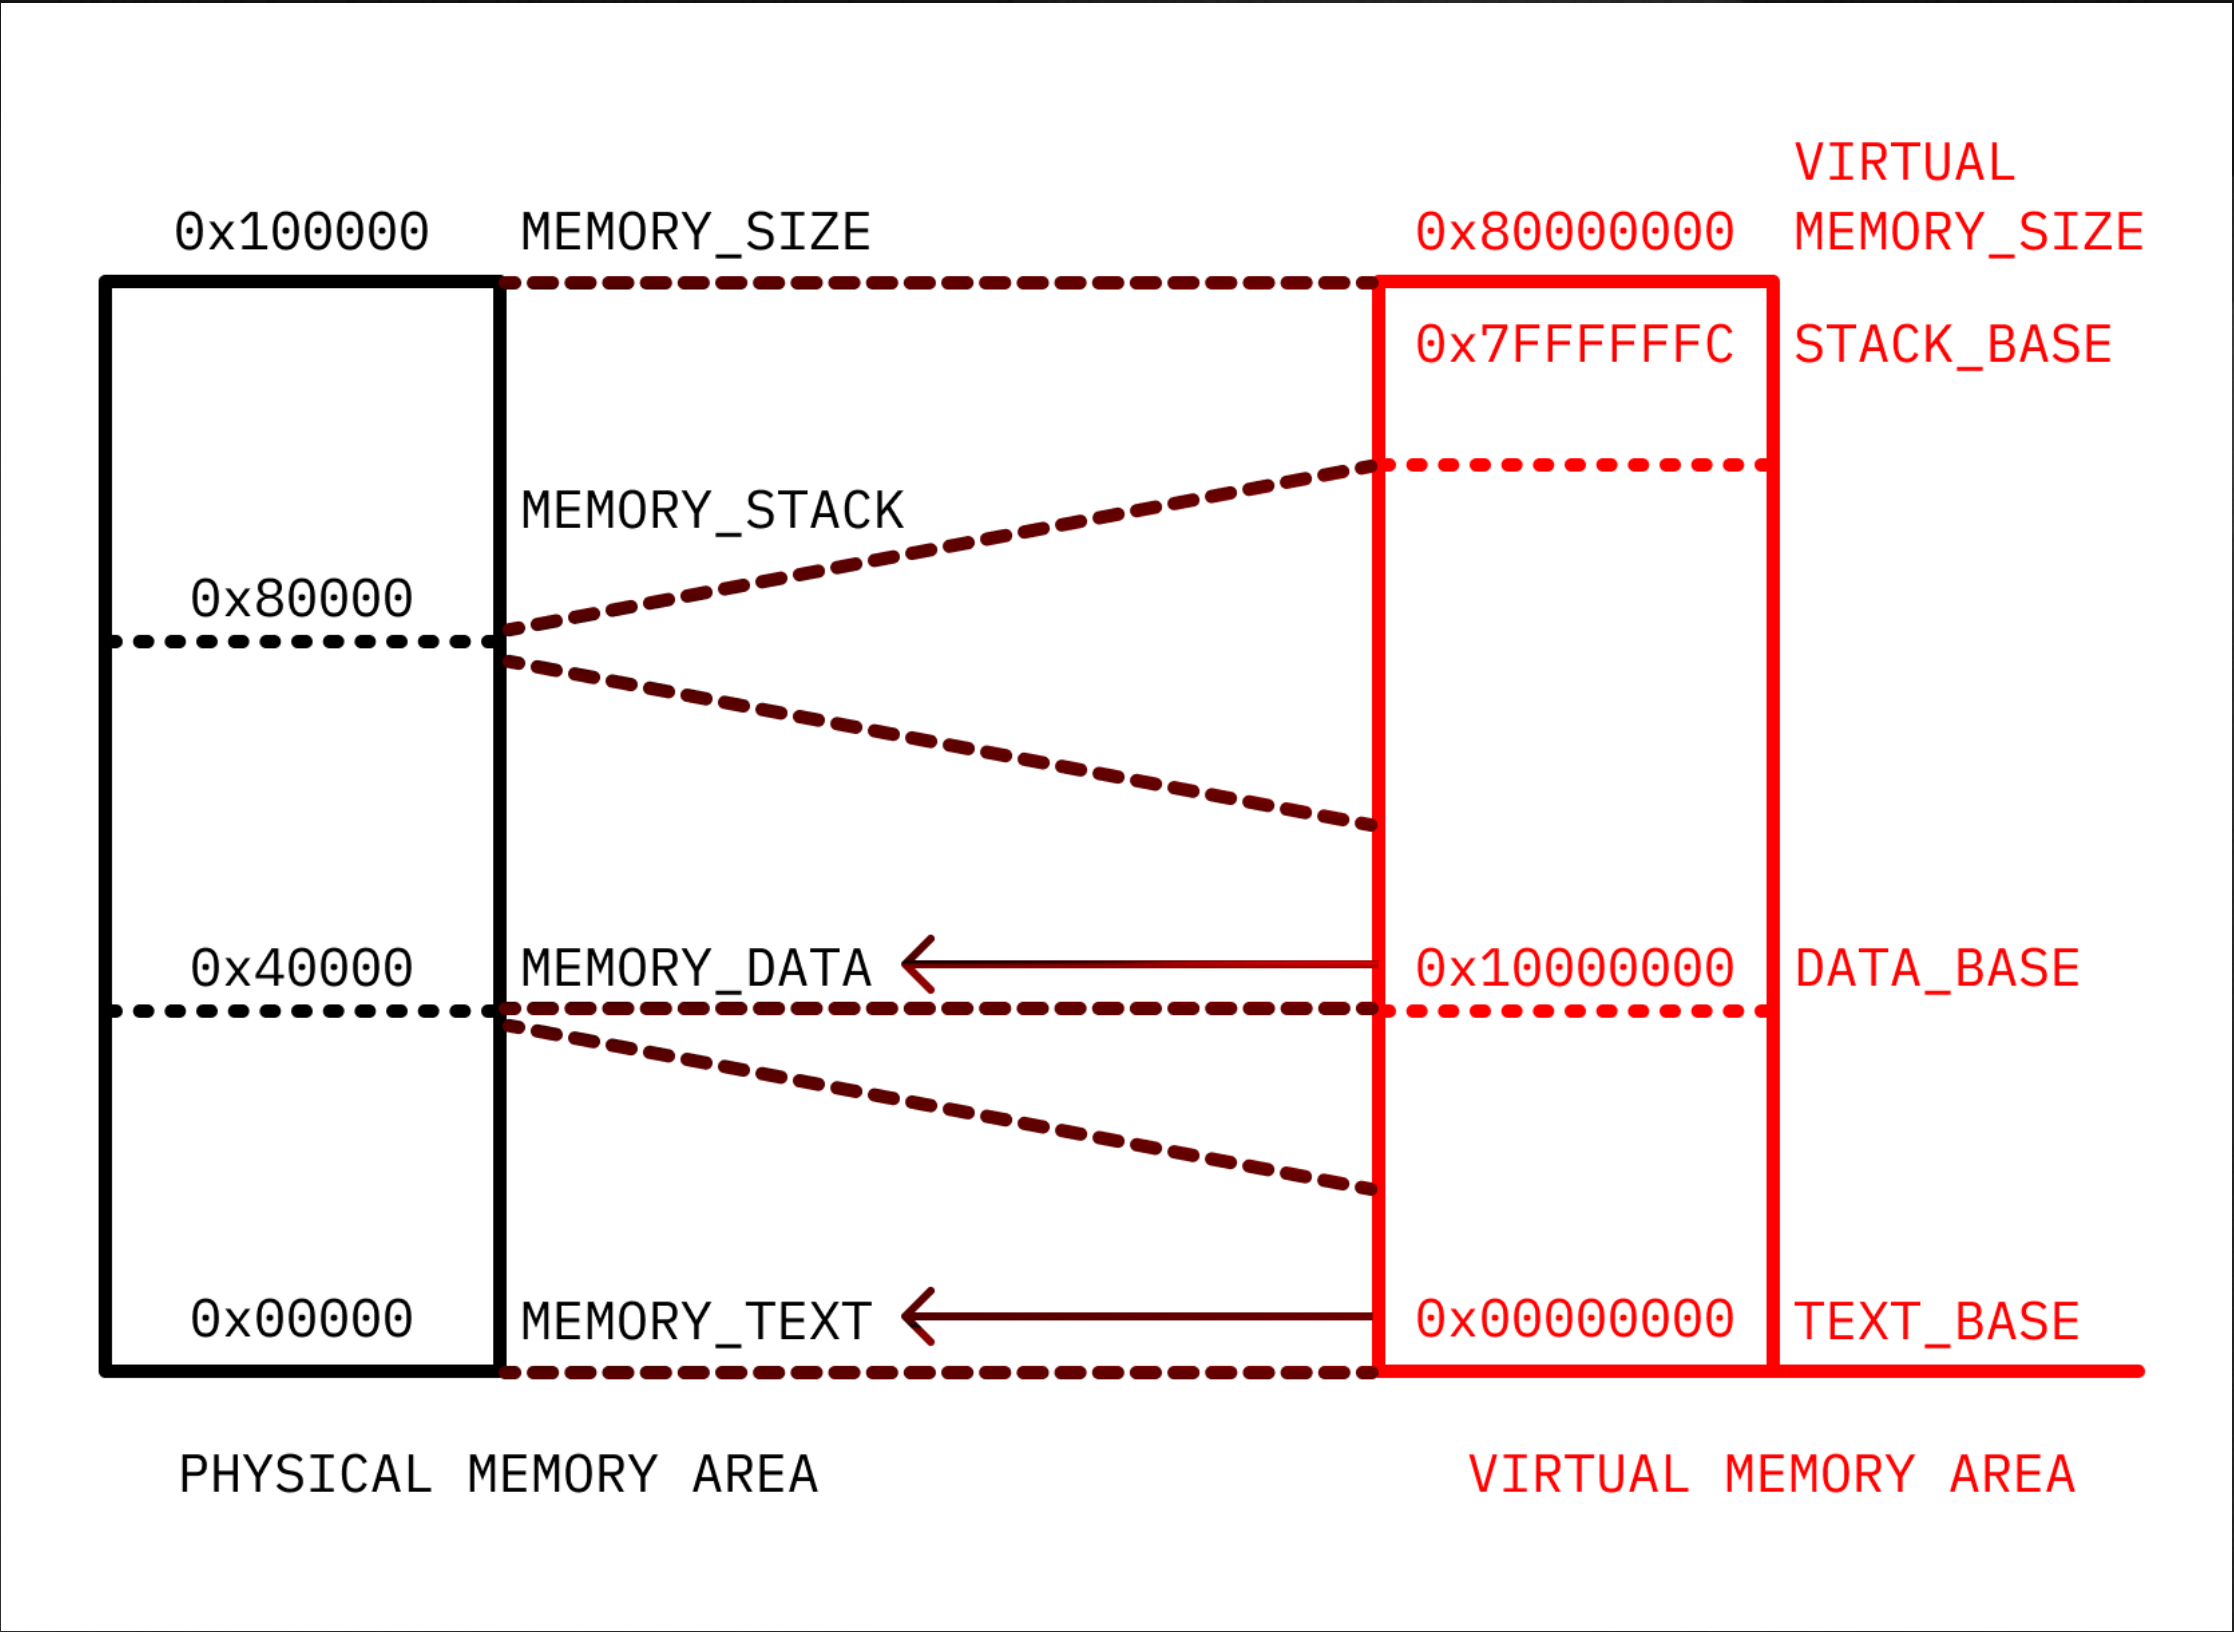
## A modern solution

- Idea: hide physical address from programs
- Introduce **Virtual Address Space**



# Virtual Memory — Base and Bound

Segment	Virtual Base	Physical Base	Bound
Text	0x00000000	0x00000	0x40000
Data	0x10000000	0x40000	0x40000
stack	0x7FF80000	0x80000	0x20000
heap	-	-	-





# Virtual Memory — Base and Bound

Memory fragmentation



$1GB - 300MB - 300MB - 300MB = 124MB$

Program	System	Prog1	Prog2			
Size	300M	300M	300M			



# Virtual Memory — Base and Bound

## Memory fragmentation



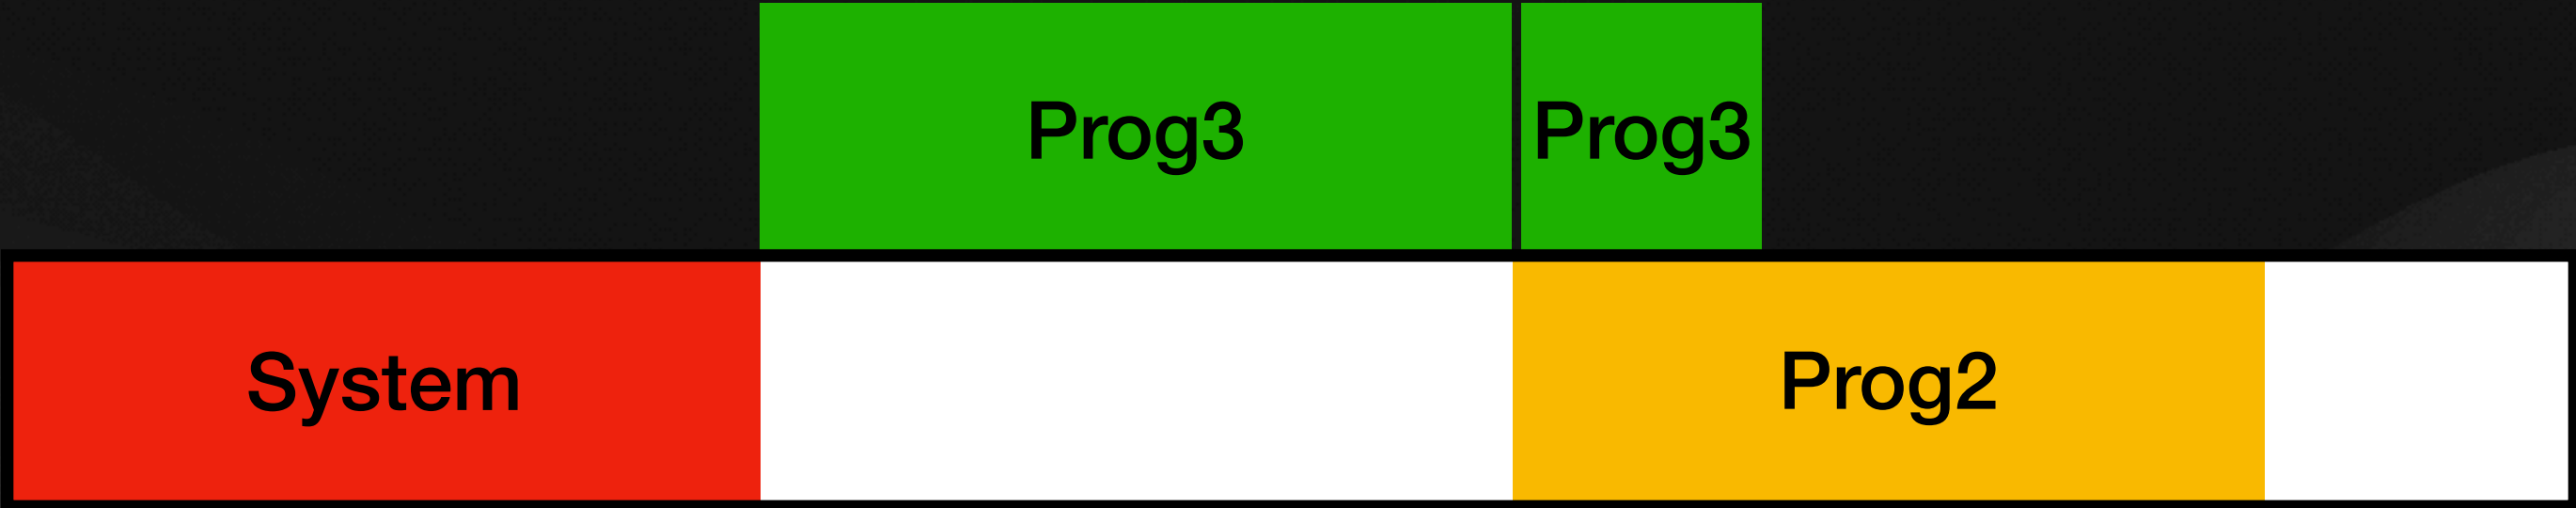
$1GB - 300MB - 300MB = 424MB$

Program	System	Prog1	Prog2			
Size	300M	-	300M			



# Virtual Memory — Base and Bound

## Memory fragmentation



Temporary solution  
(partition)

$$1GB - 300MB - 300MB - 400MB = 24MB$$

Program	System	Prog1	Prog2	Prog3		
Size	300M	-	300M	400M		





# Why not **partition** the whole memory?

## Introducing Pages

# Pages



# Pages


## Basic idea

- See memory in pages (blocks) instead of bytes
- When a process need more memory, simply allocate one page
  - Avoid fragmentation



# Pages

## On my machine




Whole memory (32GB)



# Pages

On my machine



Whole memory (32GB)




A page (4096 Bytes)



# Pages

On my machine



Whole memory (32GB)



getconf PAGE\_SIZE

A page (4096 Bytes)



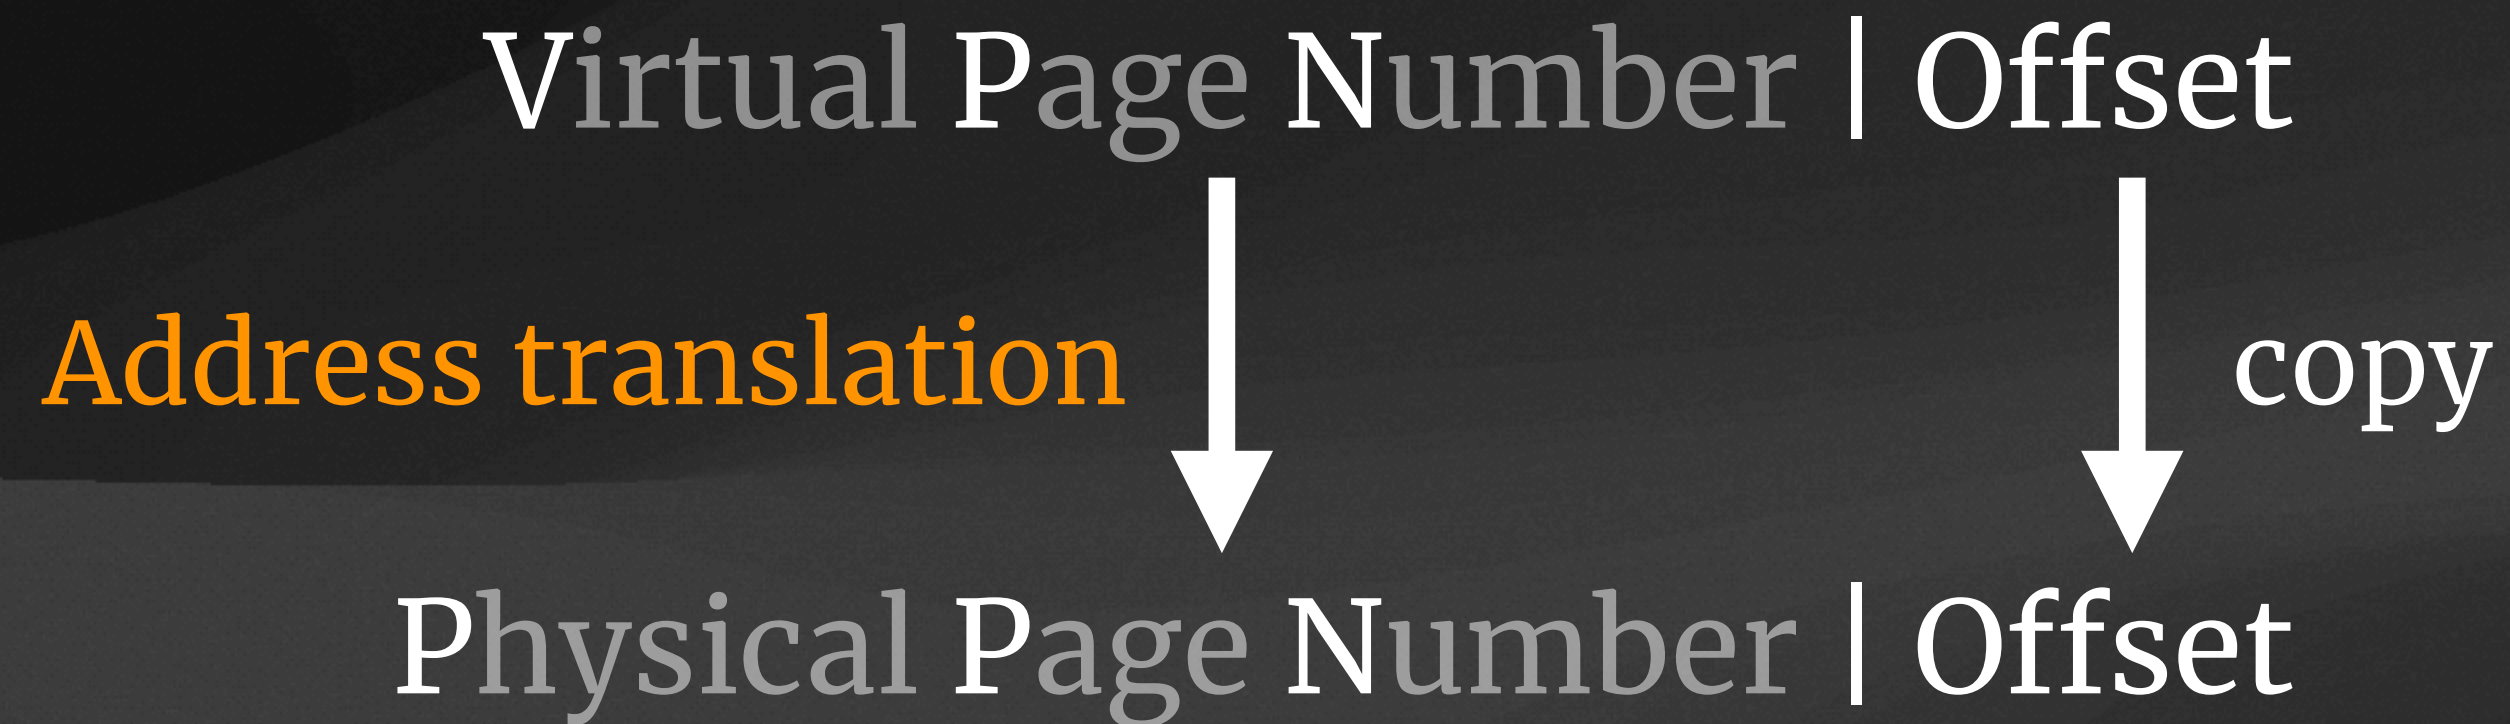
cat /proc/cpuinfo | grep cache\_alignment

A cache block (64 Bytes)



# Virtual Memory — Pages

Harness pages in **address translation**




Virtual	Physical
0b00000	0b0110
0b00001	0b0100
0b00010	-
0b00011	-
0b00100	0b0001
...	...
0b11111	0b0000



# Virtual Memory — Pages

Harness pages in **address translation**



**Simple Page Table**

	Physical
0	0b0110
1	0b0100
2	-
3	-
4	0b0001
...	...
31	0b0000

Page Table Entry →



# Virtual Memory — Pages

## Why **NOT** Simple Page Table?

32-bit Address, 4096 Bytes page, 4 Bytes PTE

- Calculate the size of simple page table:

1. *# of offset bits*  $= \log_2 4096 = 12$

2. *# of bits of virtual page number*  $= 32 - 12 = 20$

3. *# of PTE*  $= 2^{20}$

4. *Size of Page Table*  $= 4 \times 2^{20} = 2^{22} \text{ Bytes} = 4\text{MB}$



# Virtual Memory — Pages

## Why **NOT** Simple Page Table?

48-bit Address, 4096 Bytes page, 4 Bytes PTE

- Calculate the size of simple page table:

1. *# of offset bits*  $= \log_2 4096 = 12$

2. *# of bits of virtual page number*  $= 48 - 12 = 36$


3. *# of PTE*  $= 2^{36}$

4. *Size of Page Table*  $= 4 \times 2^{36} = 2^{38} \text{ Bytes} = 256\text{GB}$



# Virtual Memory — Pages

We add hierarchy to Page Table







# Virtual Memory — Pages

We add hierarchy to Page Table

Address translation





# Virtual Memory — Pages

Let's do some easy calculation!

4 Bytes PTE

- Given length of each segment — VPN1: 10 | VPN2: 10 | Offset: 12

1. *Size of page =  $2^{12}$  Bytes = 4096 Bytes*

2. *# of PTE in a L2 page table =  $2^{10}$*

3. *Size of a L2 page table =  $2^{10} \times 4 = 2^{12} = 4096$  Bytes*



# Exercise

- Physical memory : 8GB
- Page size : 8KB
- Virtual address : 46-bit
- PTE : 4B
- Assume every page table exactly fits into a single page.
- How many levels of page tables would be required.



# Virtual Memory — Translation Lookaside Buffer

Recall: Memory & (fully associative) Cache

	Data
0	0x34
1	0x7c
2	0x19
3	0x01
4	0x99
...	...
31	0x12

Tag	Data
01101101	0x91e989af
01000010	0xf4958104
11111011	0x82a0f265
01100101	0xfcf4ceef
00010101	0x88b5af65
10110101	0x190349ff
10011010	0x4358f247



# Virtual Memory — Translation Lookaside Buffer

## (Simple) page table & TLB

	Physical
0	0x34
1	0x7c
2	0x19
3	0x01
4	0x99
...	...
31	0x12

Virtual	Physical
0b10010	0b01110
0b11101	0b01100
0b00010	0b00001
0b00011	0b10111
0b01100	0b00001
0b11111	0b11101
0b10011	0b01100



# Virtual Memory — Translation Lookaside Buffer

- Fully associative most of the time
- Can also apply replacement policy
- Each program has different address space, so invalidate TLB when switching
- TLB reach: Assume simple page table, 64 entries TLB, 4KB page
  - $TLB\ reach = 64 \times 4\ KB = 256\ KB$