



信息科学与技术学院

School of Information Science and Technology

CS 110

Computer Architecture

RISC-V

Instructors:

Siting Liu & Chundong Wang

Course website: [https://toast-](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2024/index.html)

[lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2024/index.html](https://toast-lab.sist.shanghaitech.edu.cn/courses/CS110@ShanghaiTech/Spring-2024/index.html)

School of Information Science and Technology (SIST)

ShanghaiTech University

2024/3/12

Administrative

- HW2, due Mar. 22nd and Lab2 (with guidance to valgrind) are released
- Start early on labs so that checking can be faster
- Proj1.1 will be released on Piazza this week
- Discussion this week on memory management & valgrind (useful for your labs/HWs/projects) by TA Suting Chen at teaching center 301 on Monday
 - Discussion on Mar. 15th will be held at SPST 4-122 (only this one)
- The similarity check is conducted automatically for each HW, proj, etc. Please comply with the course rules!
- We grade only on your most recent submissions before ddl for all assignments!

Outline

- Intro. to ISA
- Intro. to RISC-V
- Assembly instructions in RISC-V (RV32I)
 - R-type
 - I-type arithmetic and logic
 - I-type load
 - S-type
 - Decision-making instructions

Intro to ISA

- Part of the **abstract model** of a computer that defines how the CPU is controlled by the software; **interface** between the hardware and the software;
- **Programmers' manual** because it is the portion of the machine that is visible to the **assembly language programmers**, the compiler writers, and the application programmers.
- Defines the supported data types, the registers, how the hardware manages main memory, key features, **instructions that can be executed (instruction set)**, and the input/output model of multiple ISA implementations
- ISA can be extended by adding instructions or other capabilities

-by ARM

ISA vs. Microarchitecture

- ISA: Manual for building microarchitecture or processors
- Microarchitecture: Implementation of an ISA

Popular ISAs

- X86/AMD64
 - Dominant architecture for personal computers and servers
 - Name derived from Intel 8086/80186/80286...
 - Multiple version: X86-16, X86-32 (IA-32), X86-64 (AMD64)
 - Extensions such as MMX, SSE, etc.
- Major vendors
 - Intel, AMD, VIA, Zhaoxin, DM&P, RDC
 - To OEM (original equipment manufacturer)
- CISC (complex instruction set computer)
 - Variable-length instructions
 - Allow memory access with instructions other than load or store
 - VAX architecture had an instruction to multiply polynomials!

add: X86 integer addition

Syntax

add <reg>, <reg>

add <reg>, <mem>

add <mem>, <reg>

add <mem>, imm*

... ..

Popular ISAs

- ARM
 - Dominant architecture for embedded devices
 - Advanced RISC Machine
 - Multiple version: ARMv1-ARMv9
- Major vendors
 - Apple, Huawei, Qualcomm, Xilinx, etc.
 - ARM sells IP cores to IC vendors (core licence)
 - IC vendors sell MCU/CPU/SoC to OEM or for self use
- RISC (reduced instruction set computer)
 - 32-bit fixed-length instructions (not actually for Thumb-16)
 - Allow memory access with only load or store instructions
 - Simpler to design hardware. Generally generate smaller heat

add: ARM addition

Syntax

add(S) <reg>, <reg>, <reg or imm>

RISC vs. CISC

Disassembly of section __TEXT,__text:

0000000000000000 <ltmp0>:

```

0: ff c3 00 d1 : sub sp, sp, #48
4: fd 7b 02 a9 : stp x29, x30, [sp, #32]
8: fd 83 00 91 : add x29, sp, #32
c: 08 00 80 52 : mov w8, #0
10: e8 0f 00 b9 : str w8, [sp, #12]
14: bf c3 1f b8 : stur wzr, [x29, #-4]
18: 48 9a 80 52 : mov w8, #1234
1c: a8 83 1f b8 : stur w8, [x29, #-8]
20: 28 1c 82 52 : mov w8, #4321
24: a8 43 1f b8 : stur w8, [x29, #-12]
28: a8 83 5f b8 : ldur w8, [x29, #-8]
2c: a9 43 5f b8 : ldur w9, [x29, #-12]
30: 08 01 09 0b : add w8, w8, w9
34: e8 13 00 b9 : str w8, [sp, #16]
38: e9 13 40 b9 : ldr w9, [sp, #16]
3c: e8 03 09 aa : mov x8, x9
40: e9 03 00 91 : mov x9, sp
44: 28 01 00 f9 : str x8, [x9]
48: 00 00 00 90 : adrp x0, 0x0 <ltmp0+0x48>
4c: 00 00 00 91 : add x0, x0, #0
50: 00 00 00 94 : bl 0x50 <ltmp0+0x50>
54: e0 0f 40 b9 : ldr w0, [sp, #12]
58: fd 7b 42 a9 : ldp x29, x30, [sp, #32]
5c: ff c3 00 91 : add sp, sp, #48
60: c0 03 5f d6 : ret

```

Assembly

Compiled on Mac machine using ARM CPU

0000000000000054 <main>:

```

54: 55 : push %rbp
55: 48 89 e5 : mov %rsp,%rbp
58: 48 83 ec 30 : sub $0x30,%rsp
5c: e8 00 00 00 00 : call 61 <main+0xd>
61: c7 45 fc d2 04 00 00 : movl $0x4d2,-0x4(%rbp)
68: c7 45 f8 e1 10 00 00 : movl $0x10e1,-0x8(%rbp)
6f: 8b 55 fc : mov -0x4(%rbp),%edx
72: 8b 45 f8 : mov -0x8(%rbp),%eax
75: 01 d0 : add %edx,%eax
77: 89 45 f4 : mov %eax,-0xc(%rbp)
7a: 8b 45 f4 : mov -0xc(%rbp),%eax
7d: 89 c2 : mov %eax,%edx
7f: 48 8d 05 00 00 00 00 : lea 0x0(%rip),%rax # 86 <main+0x32>
86: 48 89 c1 : mov %rax,%rcx
89: e8 72 ff ff ff : call 0 <printf>
8e: b8 00 00 00 00 : mov $0x0,%eax
93: 48 83 c4 30 : add $0x30,%rsp
97: 5d : pop %rbp
98: c3 : ret
99: 90 : nop
9a: 90 : nop
9b: 90 : nop
9c: 90 : nop
9d: 90 : nop
9e: 90 : nop
9f: 90 : nop

```

Assembly

Compiled on Windows machine using Intel CPU

Popular ISAs

- RISC philosophy (John Cocke IBM, John Hennessy Stanford, David Patterson Berkeley, 1980s)
- Hennessy & Patterson won ACM A.M. Turing Award (2017)

Reduced Instruction Set Computer (RISC)

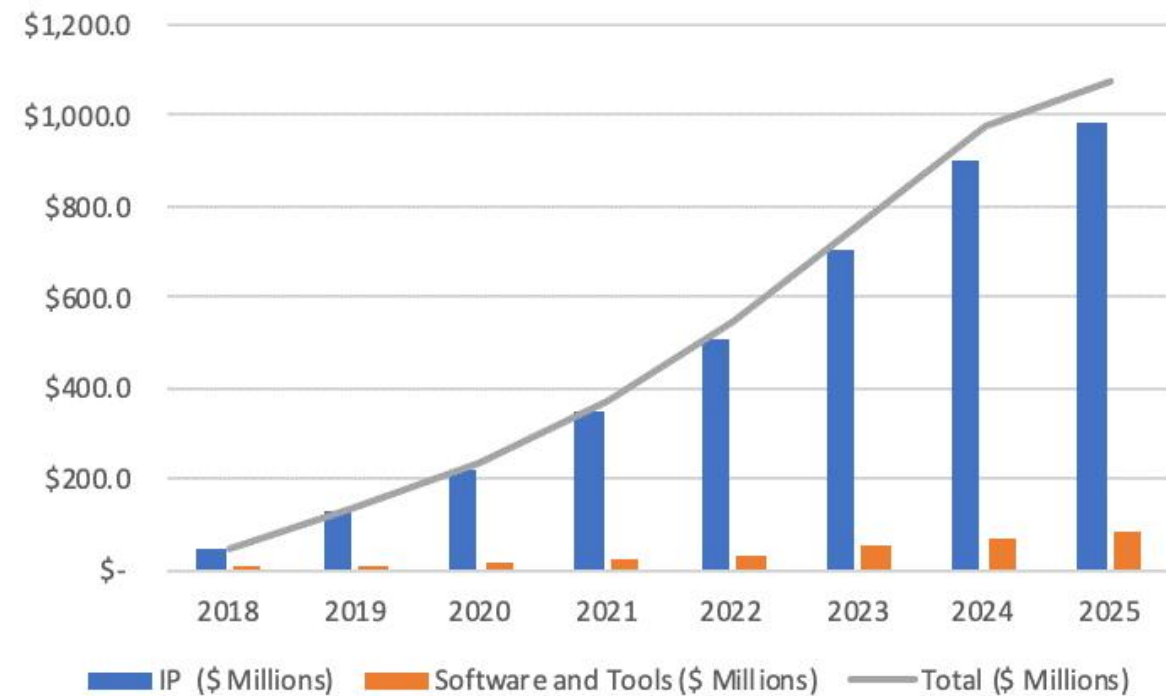
- Keep the instruction set small and simple, makes it easier to build fast hardware.
- Let software do complicated operations by composing simpler ones.

Popular ISAs

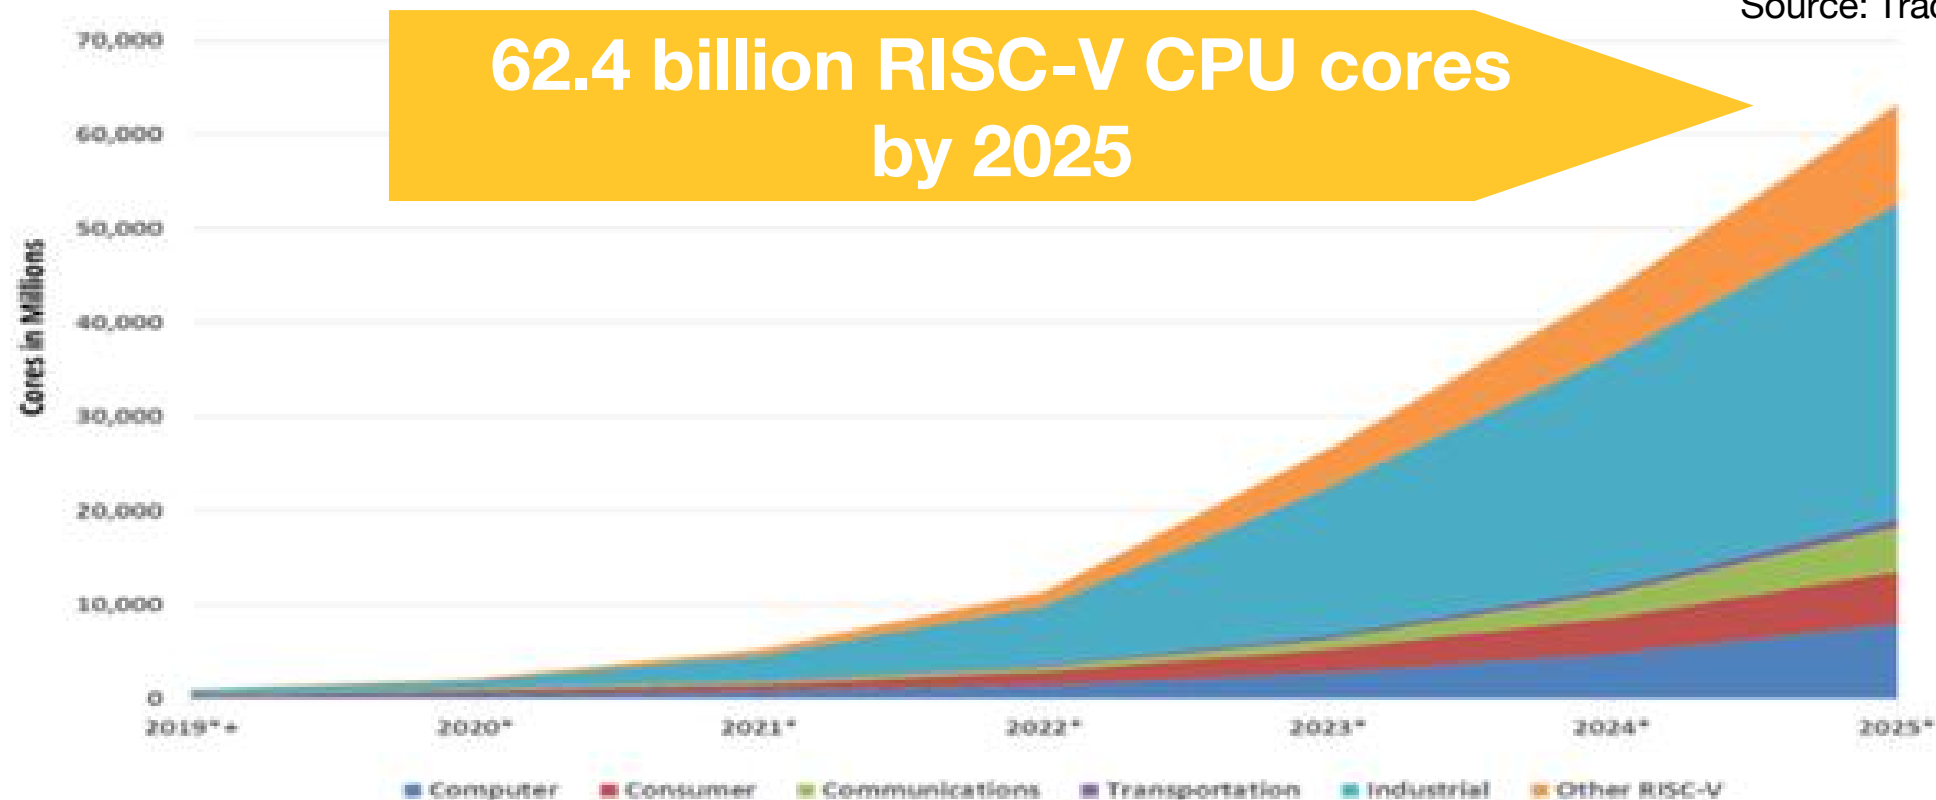
- RISC-V (pronounced “risk-five”)
 - Started as a summer project in UC Berkeley, 2010
 - The ISA itself is published in 2011 as open source
 - RISC-V foundation formed 2015 to own, maintain and publish IP related to RISC-V’s definition (a nonprofit business association)
- More than 3,100 members and still growing
 - Alibaba Cloud: T-Head 玄铁 C series; E series, and R series
 - Huawei: Hi3861V100 SoC for IoT/smart home
 - Tencent: a premier member
 - Intel, Google, Meta, SiFive, AMD/Xilinx, etc.
 - ShanghaiTech hold several RISC-V Summits China recent years!
- Other ISA examples: MIPS, IBM/Motorola PowerPC (quite old Mac), Intel IA64, ...

More than 3,100 RISC-V Members

The total market for RISC-V IP and Software is expected to grow to \$1.07 billion by 2025 at a CAGR of 54.1%



Source: Tractica



RISC-V

- Why RISC-V instead of Intel x86?
 - RISC-V is simple, elegant and open-source. Don't want to get bogged down in gritty details.
- It is flexible
 - Enabled by different extensions

Name	Description
Base	
RV32I	Base Integer Instruction Set, 32-bit
RV32E	Base Integer Instruction Set (embedded), 32-bit, 16 register
RV64I	Base Integer Instruction Set, 64-bit
RV64E	Base Integer Instruction Set (embedded), 64-bit
RV128I	Base Integer Instruction Set, 128-bit
Extension	
M	Standard Extension for Integer Multiplication and Division
A	Standard Extension for Atomic Instructions
F	Standard Extension for Single-Precision Floating-Point
D	Standard Extension for Double-Precision Floating-Point

$$RVG = RVI + M + A + F + D$$

Manual available: <https://riscv.org>

Where are we?

High Level Language
Program (e.g., C)

Compiler

Assembly Language
Program (e.g., RISC-V)

Assembler

Machine Language
Program (RISC-V)

```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

```
lw    t0, 0(s2)
lw    t1, 4(s2)
sw    t1, 0(s2)
sw    t0, 4(s2)
```

We are here!

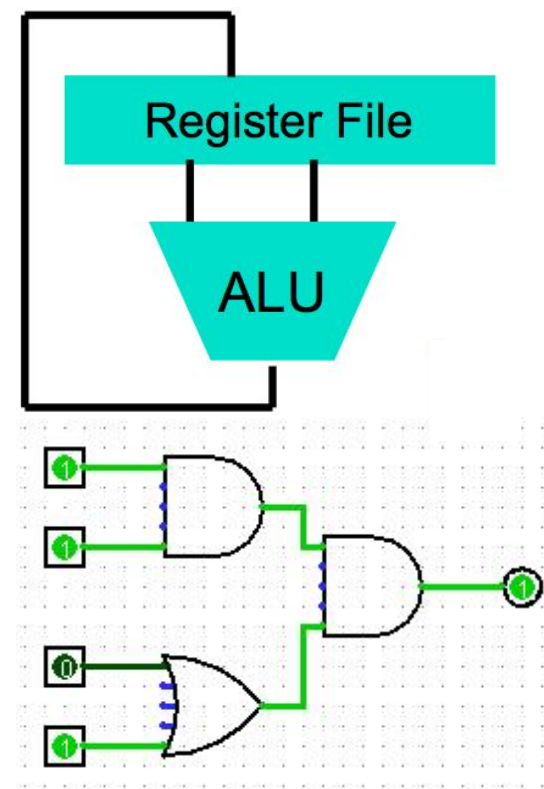
```
0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111
```

Machine
Interpretation

Hardware Architecture Description
(e.g., block diagrams)

Architecture
Implementation

Logic Circuit Description
(Circuit Schematic Diagrams)



Assembly Language

- Basic job of a CPU: execute a series of instructions.
- Instructions are the primitive operations that the CPU may execute.
- Basic job of a instruction: change the state of a computer.

Disassembly of section __TEXT,__text:

```
0000000000000000 <tmp0>:
 0: ff c3 00 d1  sub sp, sp, #48
 4: fd 7b 02 a9  stp x29, x30, [sp, #32]
 8: fd 83 00 91  add x29, sp, #32
 c: 08 00 80 52  mov w8, #0
10: e8 0f 00 b9  str w8, [sp, #12]
14: bf c3 1f b8  stur wzr, [x29, #-4]
18: 48 9a 80 52  mov w8, #1234
1c: a8 83 1f b8  stur w8, [x29, #-8]
20: 28 1c 82 52  mov w8, #4321
24: a8 43 1f b8  stur w8, [x29, #-12]
28: a8 83 5f b8  ldur w8, [x29, #-8]
2c: a9 43 5f b8  ldur w9, [x29, #-12]
30: 08 01 09 0b  add w8, w8, w9
34: e8 13 00 b9  str w8, [sp, #16]
38: e9 13 40 b9  ldr w9, [sp, #16]
3c: e8 03 09 aa  mov x8, x9
40: e9 03 00 91  mov x9, sp
44: 28 01 00 f9  str x8, [x9]
48: 00 00 00 90  adrp x0, 0x0 <tmp0+0x48>
4c: 00 00 00 91  add x0, x0, #0
50: 00 00 00 94  bl 0x50 <tmp0+0x50>
54: e0 0f 40 b9  ldr w0, [sp, #12]
58: fd 7b 42 a9  ldp x29, x30, [sp, #32]
5c: ff c3 00 91  add sp, sp, #48
60: c0 03 5f d6  ret
```

Exercise 1: Can we execute this assembly on a X86 CPU?

Exercise 2: Can we use a Linux OS to run ARM assembly?

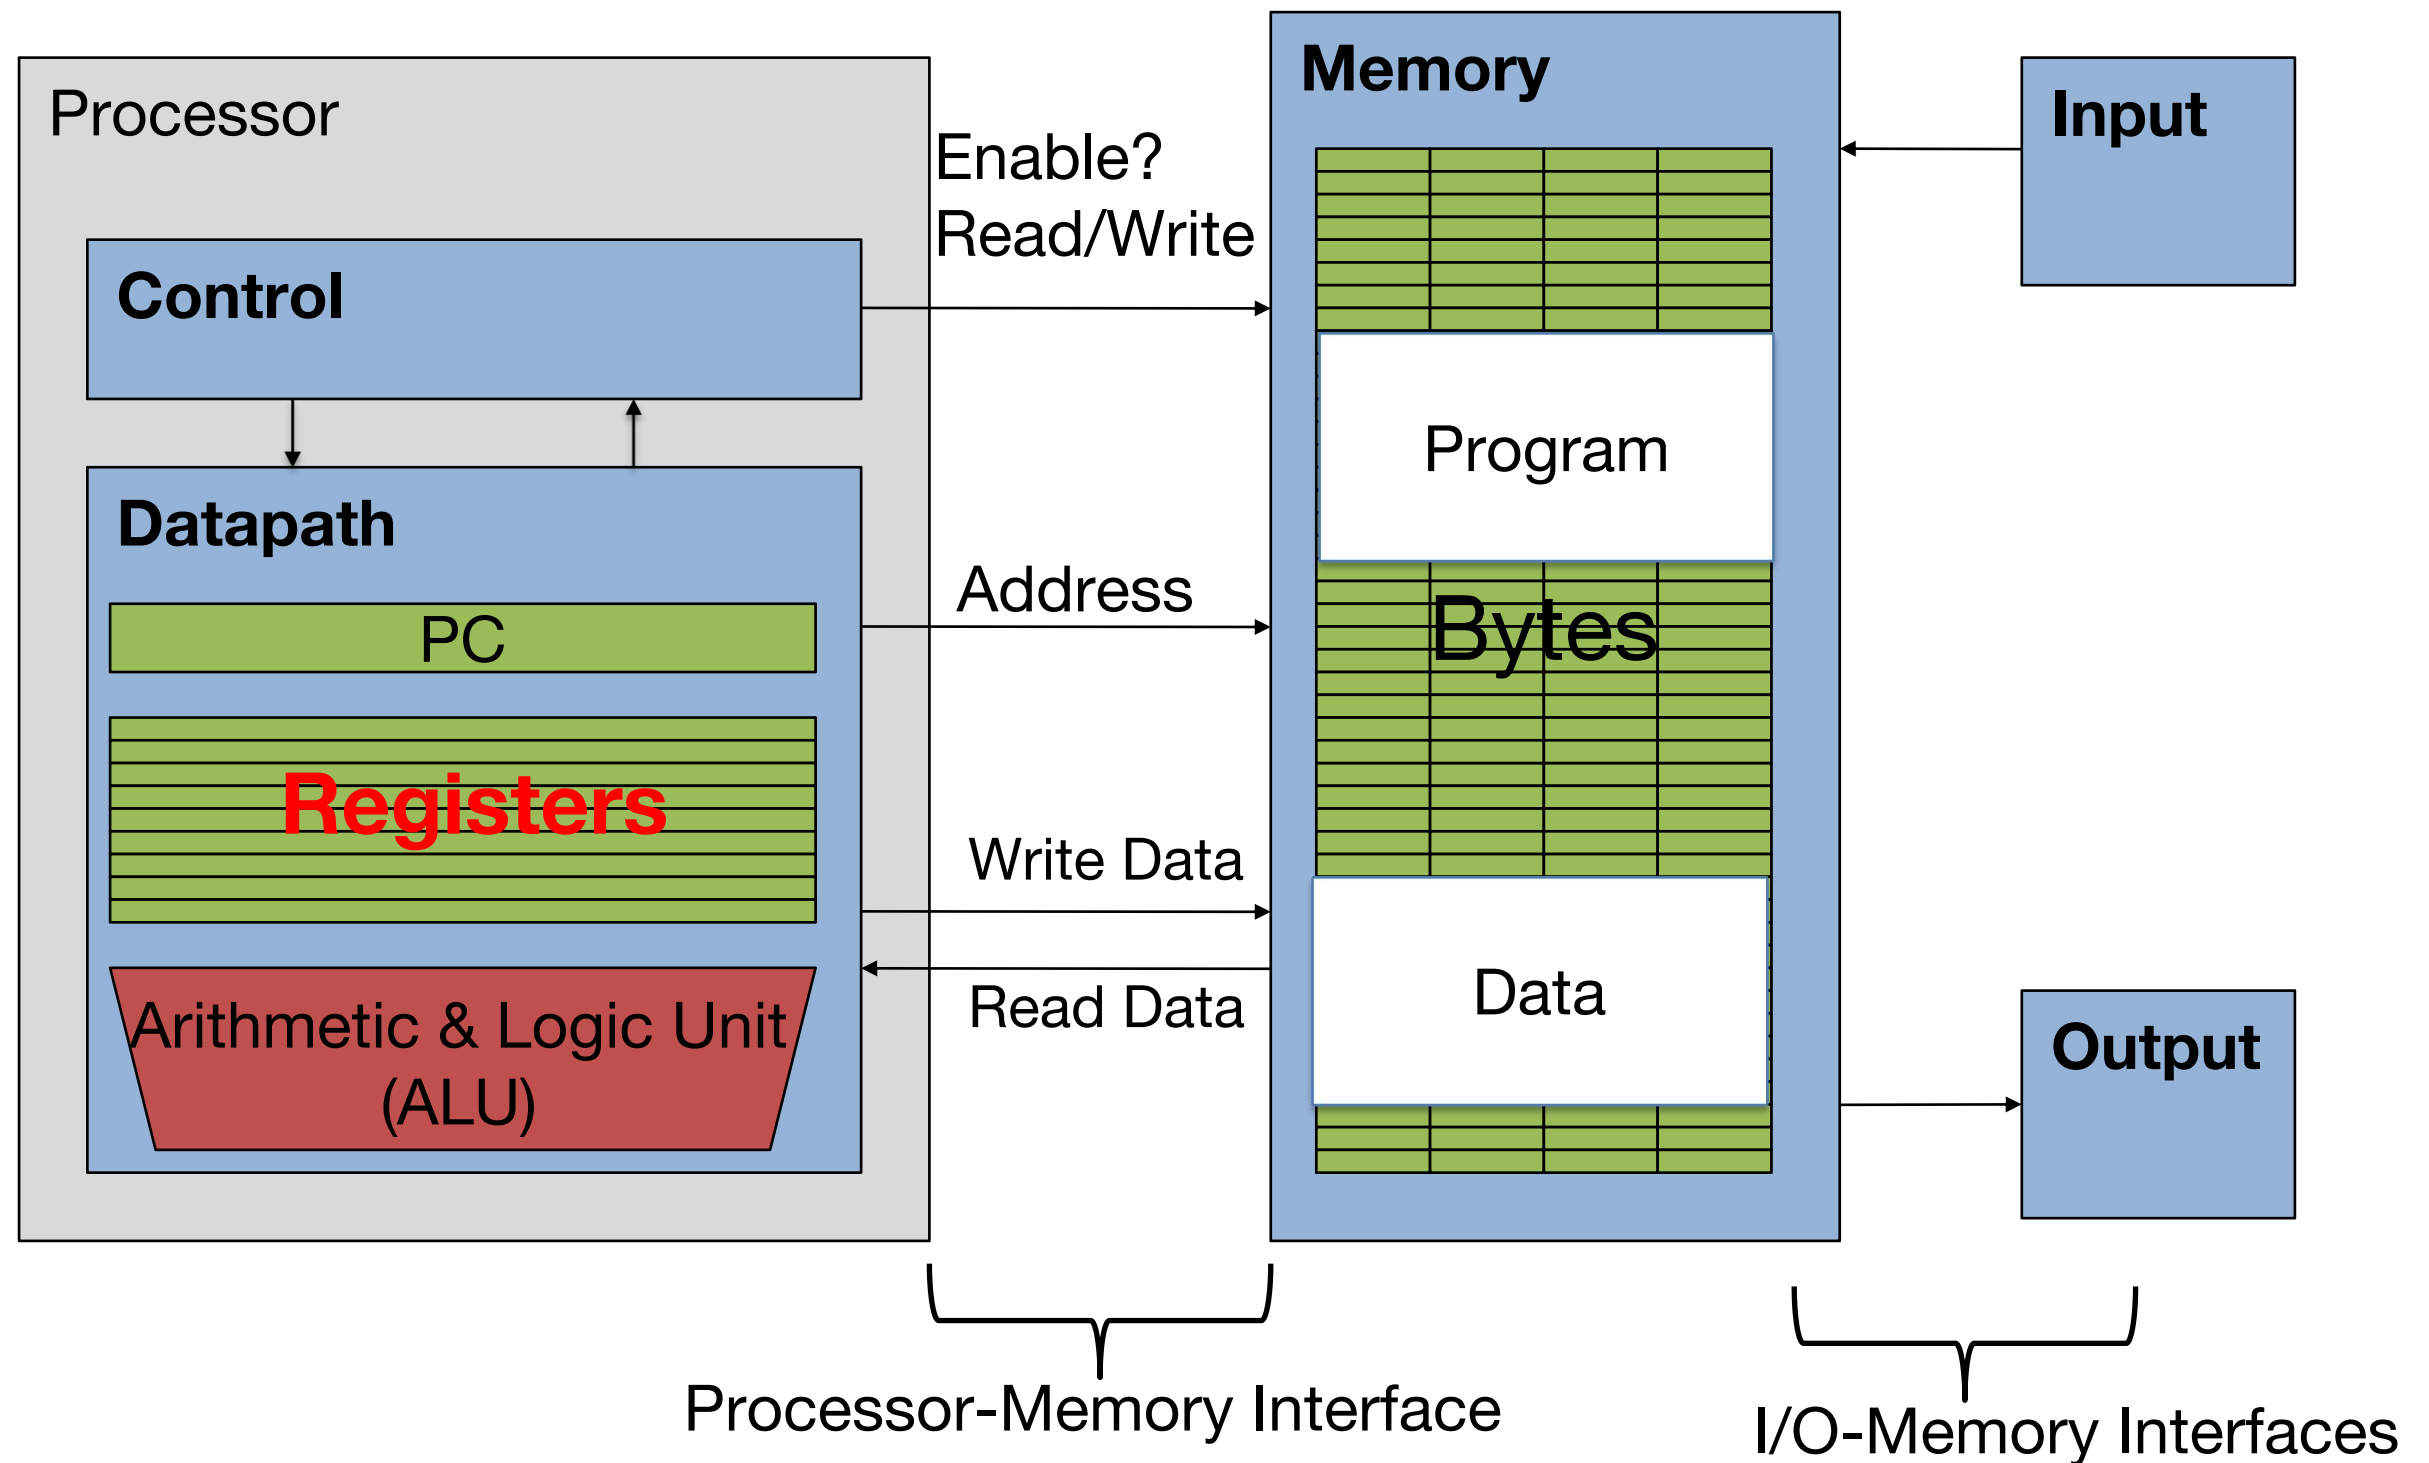
ARM Assembly

Compiled on Mac machine using ARM CPU

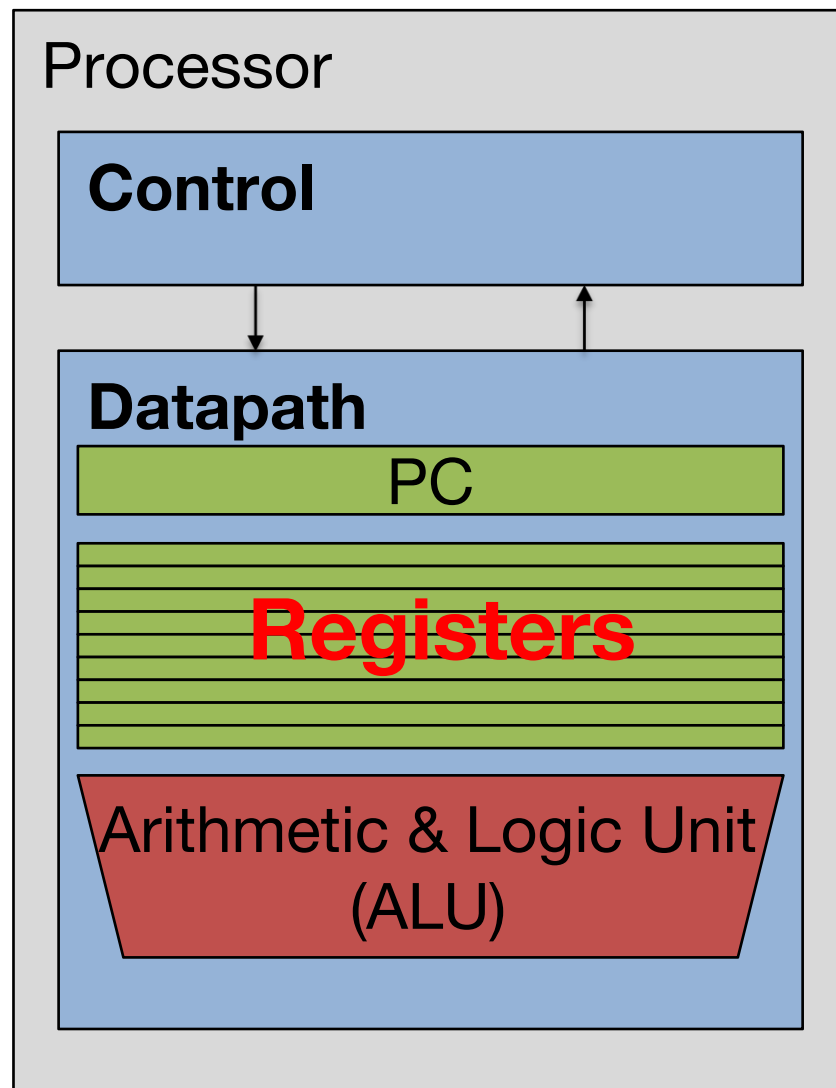
CPU State: Assembly Registers (hardware/variable)

- Unlike C or Java, assembly cannot use variables
 - Keep assembly/computer hardware abstract simple
- Assembly operands are registers
 - Limited number of special locations/memory built directly into the CPU
 - Operations can only be performed on these registers in RISC-V
- Benefit: Since registers are directly in hardware (CPU), they are very fast

Registers, inside the Processor



RV32I Registers



Registers



- Similar to memory, use “address” to refer to specific location

PC register

- Hold address of the current instruction

- 32 registers in RISC-V
 - Why 32? Smaller is faster, but too small is bad.

```

1c: a8 83 1f b8  stur w8, [x29, #-8]
20: 28 1c 82 52  mov w8, #4321
24: a8 43 1f b8  stur w8, [x29, #-12]
28: a8 83 5f b8  ldur w8, [x29, #-8]
2c: a9 43 5f b8  ldur w9, [x29, #-12]
30: 08 01 09 0b  add w8, w8, w9

```

- Each RV32 register is 32-bit wide
 - Groups of 32 bits called a word in RV32; P&H textbook uses 64-bit variant RV64 (doubleword)

C, Java variables vs. registers

- In C (and most high level languages) variables declared first and given a type
 - Example: `int fahr, celsius;`
`char a, b, c, d, e;`
- Each variable can ONLY represent a value of the type it was declared as (cannot mix and match int and char variables).
- In Assembly Language, registers have no type, **simply stores 0s and 1s**; operation determines how register contents are treated (think about the hardware)

Assembly Instructions

- In assembly language, each statement (called an instruction), executes exactly one of a short list of simple commands
- Unlike in C (and most other high level languages), each line of assembly code contains at most 1 instruction
- Another way to make your code more readable: comments!
- Hash (#) is used for RISC-V comments
 - anything from hash mark to end of line is a comment and will be ignored

```
1c: a8 83 1f b8    stur w8, [x29, #-8]
20: 28 1c 82 52    mov w8, #4321
24: a8 43 1f b8    stur w8, [x29, #-12]
28: a8 83 5f b8    ldur w8, [x29, #-8]
2c: a9 43 5f b8    ldur w9, [x29, #-12]
30: 08 01 09 0b    add w8, w8, w9
```

Assembly Instructions

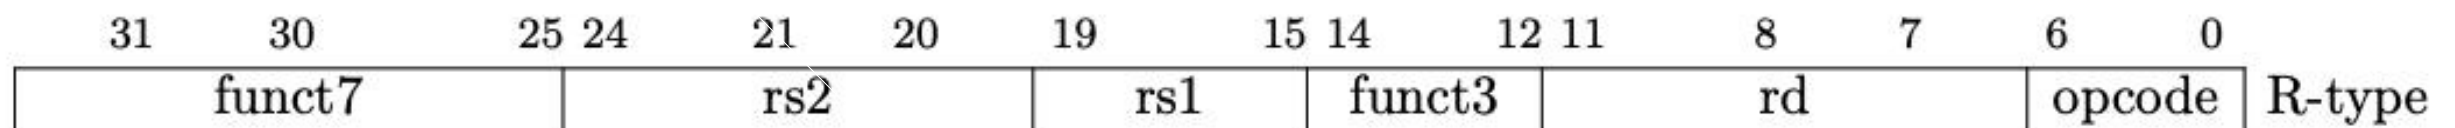
- Different types of instructions (4 core types + B/J based on the handling of immediate)

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0	
funct7				rs2			rs1		funct3		rd			opcode	R-type
imm[11:0]						rs1		funct3		rd			opcode	I-type	
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode	S-type
imm[12]	imm[10:5]			rs2			rs1		funct3		imm[4:1]	imm[11]	opcode	B-type	
imm[31:12]										rd			opcode	U-type	
imm[20]	imm[10:1]				imm[11]		imm[19:12]				rd			opcode	J-type

- Different types have different format but “rs1”, “rs2” and “rd” are at the same position (hardware friendly)
- As an ID number, the machine code of the instructions has different fields; format depends on their operands/type

Assembly Instructions

- Different types of instructions (4 core types + B/J based on the handling of immediate)



- R-type
 - Register-register operation, mainly for arithmetic & logic
 - Has two operands (accessed from the source registers, rs1 & rs2) and one output (saved to the destination register, rd)
 - Cannot access main memory (instruction executed by CPU alone, no data exchange with main memory)

RV32I R-type Arithmetic

- Syntax of instructions

- Addition: `add rd,rs1,rs2` (operation `rd,rs1,rs2`)

Adds the value stored in register `rs1` to that of `rs2` and stores the sum into register `rd`, similar to `a = b+c`, `a` \Leftrightarrow `rd`, `b` \Leftrightarrow `rs1`, `c` \Leftrightarrow `rs2`

- Example: `add x5, x2, x1`
`add x6, x0, x5`
`add x4, x1, x3`

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x1233FFFF	x4
0x12346789	x5
0x12346789	x6
	x7

$$\begin{array}{r} 12340000 \\ + 00006789 \\ \hline 0x12346789 \end{array}$$

$$\begin{array}{r} 0x12340000 \\ + 0x1233FFFF \\ \hline 1233FFFF \end{array}$$

RV32I R-type Arithmetic

- Syntax of instructions
 - Subtraction: `sub rd, rs1, rs2`

Subtract the value stored in register rs2 from that of rs1 and stores the difference into register rd, equivalent to $a = b - c$, $a \Leftrightarrow \text{rd}$, $b \Leftrightarrow \text{rs1}$, $c \Leftrightarrow \text{rs2}$

- Example: `sub x5, x2, x1`
`sub x6, x5, x0`

Handwritten calculation:

$$\begin{array}{r} 00006789 \\ - 12340000 \\ \hline \text{edcc6789} \end{array}$$

Handwritten calculation:

$$\begin{array}{l} 0x\text{edcc6789} - 0 \\ = 0x\text{edcc6789} \end{array}$$

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
	x4
0xedcc6789	x5
0xedcc6789	x6
	x7

RV32I R-type Logic Operation

- Syntax of instructions:

- AND/OR/XOR: `and/or/xor rd, rs1, rs2`

Logically **bit-wise** and/or/xor the value stored in register rs1 and that of rs2 and stores the result into register rd, equivalent to $a = b \text{ (&/|/^)}$

c, $a \Leftrightarrow \text{rd}$, $b \Leftrightarrow \text{rs1}$, $c \Leftrightarrow \text{rs2}$ *u*

- Example: `and x5, x2, x1`
`xor x6, x1, x5` *all zero*
`and x4, x1, x3` *all one*
 $x4 = x1$

*$0x12340000 \wedge 0x00000000$
 $= 0x12340000$*

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x12340000	x4
0x00000000	x5 (0x0)
0x12340000	x6
	x7

RV32I R-type Logic Operation

- Syntax of instructions:

- AND/OR/XOR: `and/or/xor rd, rs1, rs2`

Logically **bit-wise** and/or/xor the value stored in register `rs1` and that of `rs2` and stores the result into register `rd`, equivalent to $a = b \ (\&/\ /\ ^) \ c$, a

\Leftrightarrow `rd`, $b \Leftrightarrow$ `rs1`, $c \Leftrightarrow$ `rs2`

- Used for bit-mask

`and x5, x7, x4` *lower bits masked*

`or x6, x7, x4`

- Used for bit-wise negation

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0xFFFF0000	x4
0x12340000	x5
	x6
0x12345678	x7

RV32I R-type Compare

- Syntax of instructions

- SLT/SLTU: `slt/sltu rd, rs1, rs2`

Compare the value stored in register rs1 and that of rs2, sets rd=1, if rs1<rs2 otherwise rd=0, equivalent to $a = b < c ? 1 : 0$, $a \Leftrightarrow$ `rd`, $b \Leftrightarrow$ `rs1`, $c \Leftrightarrow$ `rs2`. Treat the numbers as signed/unsigned with `slt/sltu`.

- Example: `slt x5, x2, x1`
`slt x4, x3, x1`
`slt u x5, x3, x1`

- Overflow detection (unsigned)

```
add x5, x3, x3
```

```
slt u x6, x5, x3
```

- Overflow detection (signed)?

```
add t0, t1, t2
```

```
slti t3, t2, 0
```

```
slt t4, t0, t1
```

```
bne t3, t4, overflow
```

$x2 < x1 \rightarrow x5 = 1$ *Signed*

$x3 < x1 \rightarrow x4 = 1$

$x3 > x1 \rightarrow x5 = 0 \rightarrow$ unsigned

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x1	x4
0x00000001	x5 or 0x1 \rightarrow 0x0
	x6
	x7

RV32I R-type Shift

- Syntax of instructions:

- Shift left/right (arithmetic): `sll/srl/sra rd, rs1, rs2`

Left/Right shifts the value stored in register `rs1` by that of `rs2`,
 equivalent to $a = b \ll / \gg c$, $a \Leftrightarrow \text{rd}$, $b \Leftrightarrow \text{rs1}$, $c \Leftrightarrow \text{rs2}$.
 (lower 5 bits)

arithmetic: sign extended.

- Example: `sll x5, x2, x4`
`srl x6, x1, x4`
`sra x7, x3, x4`

000067890

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3 0b111...11
0x4	x4
0x000067890	x5 7111...11
0x0124000	x6
0x7F--F	x7

32 bit.

27

RV32I R-type Shift

- Syntax of instructions:

- Shift left/right (arithmetic): `sll/srl/sra rd, rs1, rs2`

Left/Right shifts the value stored in register `rs1` by that of `rs2`, equivalent to $a = b \ll / \gg c$, $a \Leftrightarrow \text{rd}$, $b \Leftrightarrow \text{rs1}$, $c \Leftrightarrow \text{rs2}$.

arithmetic: sign extended.

- Example: `sll x5, x2, x4`
`srl x6, x1, x4`
`sra x7, x3, x4`

- What is the arithmetic effect by shifting?

Registers

0	x0/zero
0x12340000	x1
0x00006789	x2
0xFFFFFFFF	x3
0x4	x4
0x00067890	x5
	x6
	x7