



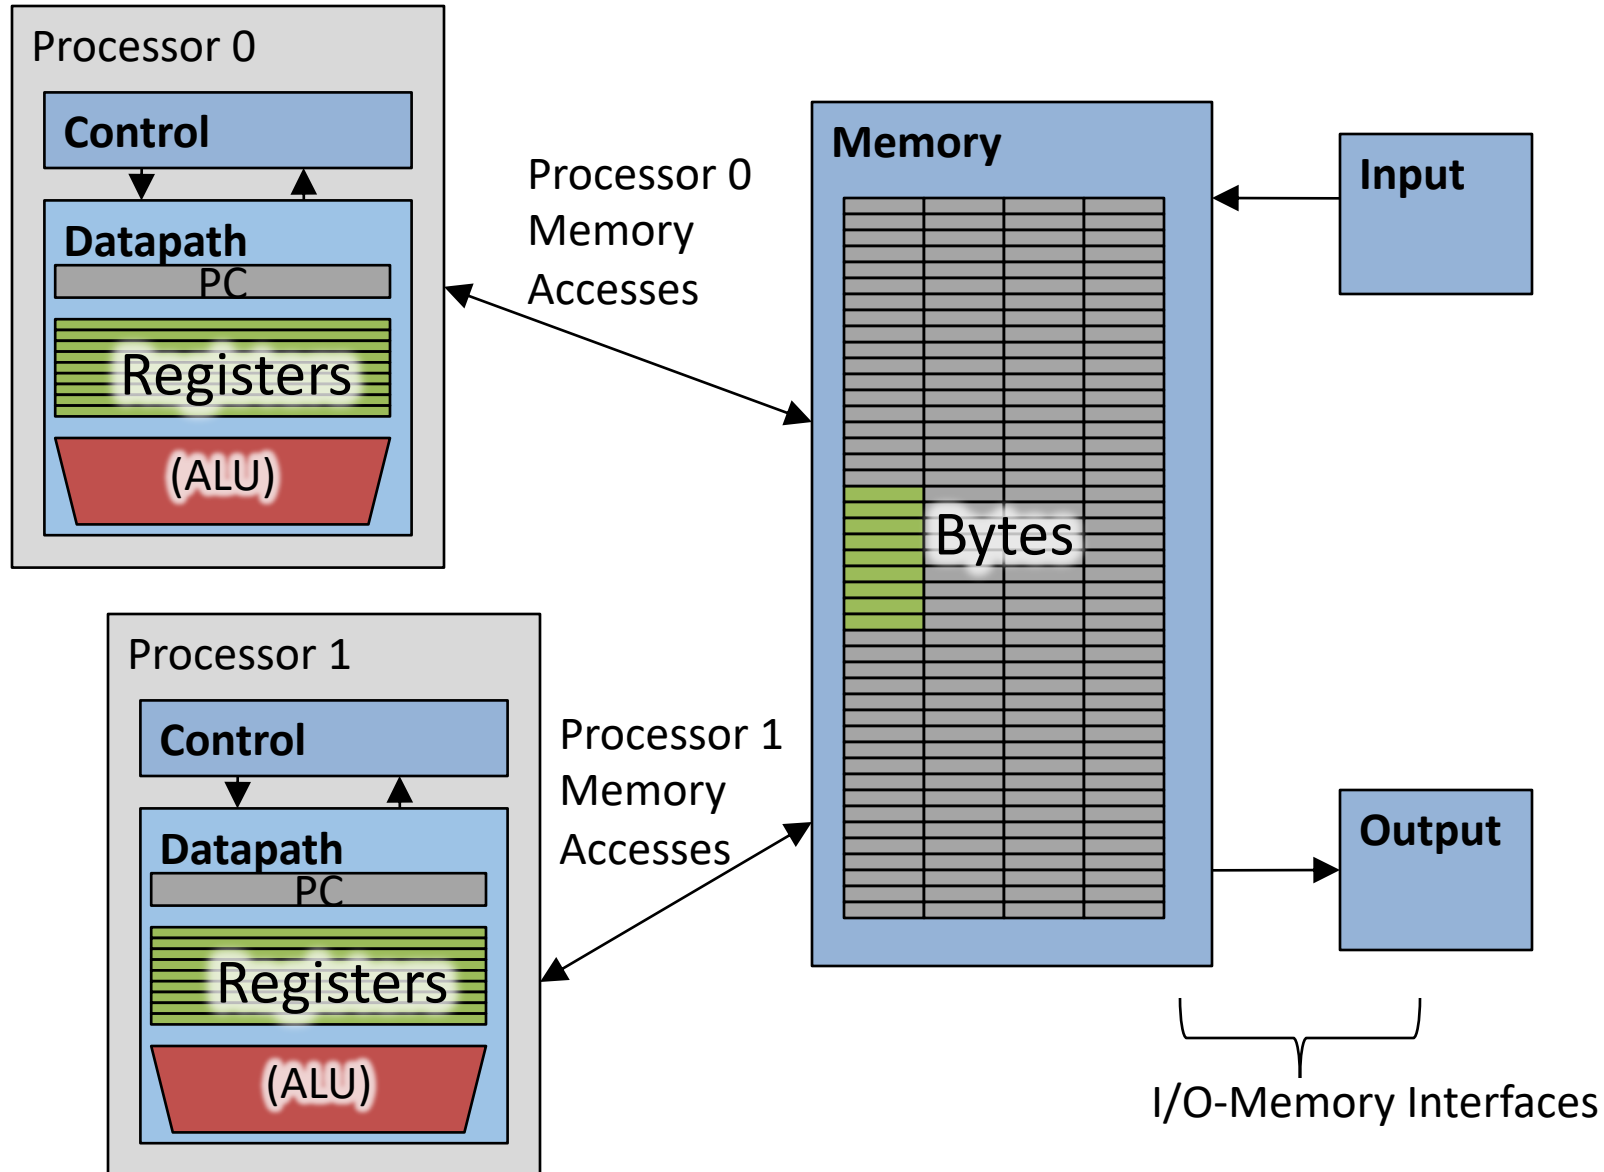
上海科技大学
ShanghaiTech University



CS110 Computer Architecture

Sync and OpenMP

Chundong Wang & Siting Liu
SIST, ShanghaiTech



Multi-core



TLP, OpenMP, and Sync

- Multicore
 - Hyperthreading
- OpenMP
 - Shared memory
 - Language extension
- Lock for synchronization
 - Data race
 - At least one write operation

```
wangc@HP:~/TT$ gcc omp.c -o p -O3 -fopenmp
wangc@HP:~/TT$ ./p
Hello World from thread = 1
Hello World from thread = 7
Hello World from thread = 6
Hello World from thread = 5
Hello World from thread = 0
Number of threads = 8
Hello World from thread = 2
Hello World from thread = 4
Hello World from thread = 3
wangc@HP:~/TT$ ./p
Hello World from thread = 3
Hello World from thread = 0
Number of threads = 8
Hello World from thread = 1
Hello World from thread = 2
Hello World from thread = 7
Hello World from thread = 5
Hello World from thread = 6
Hello World from thread = 4
wangc@HP:~/TT$ ./p
Hello World from thread = 2
Hello World from thread = 7
Hello World from thread = 5
Hello World from thread = 3
Hello World from thread = 0
Number of threads = 8
Hello World from thread = 6
Hello World from thread = 4
Hello World from thread = 1
wangc@HP:~/TT$
```



Lock Synchronization (1/2)

- Use a “Lock” to grant access to a region (*critical section*) so that only one thread can operate at a time
 - Need all processors to be able to access the lock, so use a location in shared memory as *the lock*
- Processors read lock and either wait (if locked) or set lock and go into critical section
 - **0** means lock is free / open / unlocked / lock off
 - **1** means lock is set / closed / locked / lock on



Lock Synchronization (2/2)

- Pseudocode:

```
Check lock
Set the lock
Critical section
(e.g. change shared variables)
Unset the lock
```

Can loop/idle here
if locked



Possible Lock Implementation

- Lock (a.k.a. busy wait)

```
Get_lock:                                # s0 -> addr of lock
      addiu t1,zero,1                     # t1 = Locked value
Loop:  lw      t0,0(s0)                   # load lock
      bne     t0,zero,Loop               # loop if locked
Lock:  sw      t1,0(s0)                   # Unlocked, so lock
```

- Unlock

```
Unlock:
      sw zero,0(s0)
```

- Any problems with this?



Possible Lock Problem

- Thread 1

```
    addiu t1,zero,1
Loop: lw  t0,0(s0)

    bne t0,zero,Loop

Lock: sw t1,0(s0)
```

- Thread 2

```
    addiu t1,zero,1
Loop: lw t0,0(s0)

    bne t0,zero,Loop

Lock: sw t1,0(s0)
```

Time

*Both threads think they have set the lock!
Exclusive access not guaranteed!*



Hardware Synchronization

- Hardware support required to prevent an interloper (another thread) from changing the value
 - *Atomic* read/write memory operation
 - No other access to the location allowed between the read and write
- How best to implement in software?
 - Single instr? Atomic swap of register \leftrightarrow memory
 - Pair of instr? One for read, one for write
- Needed even on uniprocessor systems
 - Interrupts can happen: can trigger thread context switches...



上海科技大学
ShanghaiTech University




RISC-V: Two solutions!

- Option 1: Read/Write Pairs
 - Pair of instructions for “linked” read and write
 - Load reserved and Store conditional
 - No other access permitted between read and write
 - Must use *shared memory* (multiprocessing)
- Option 2: Atomic Memory Operations
 - Atomic swap of register \leftrightarrow memory



Read/Write Pairs

- Load reserved: **lr rd, rs**
 - Load the word pointed to by **rs** into **rd**, and add a reservation 
- Store conditional: **sc rd, rs1, rs2**
 - Store the value in **rs2** into the memory location pointed to by **rs1**, only if the reservation is still valid and set the status in **rd**
 - Returns 0 (success) if location has not changed since the **lr**
 - Returns nonzero (failure) if location has changed:
Actual store will not take place



Synchronization in RISC-V Example

- Atomic swap (to test/set lock variable)
- Exchange contents of register and memory:
 $s4 \leftrightarrow \text{Mem}(s1)$

try:

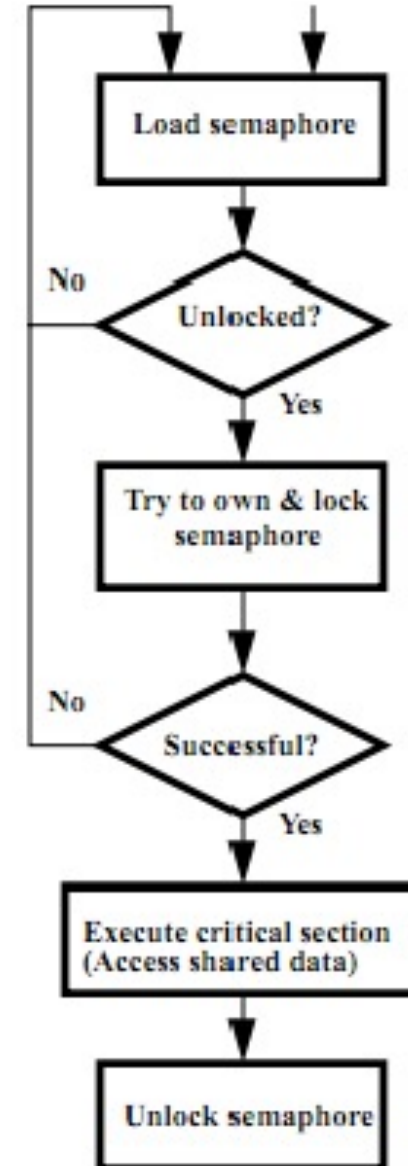
```
lr    t1, s1           #load reserved
sc    t0, s1, s4        #store conditional
bne   t0, x0, try       #loop if sc fails
add   s4, x0, t1        #load value in s4
```


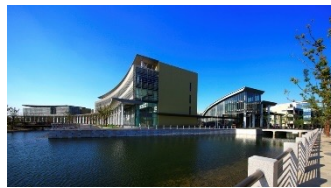
sc would fail if another threads executes **sc** here



Test-and-Set

- In a single atomic operation:
 - **Test** to see if a memory location is set (contains a 1)
 - **Set** it (to 1) if it isn't (it contained a zero when tested)
 - Otherwise indicate that the Set failed, so the program can try again
 - While accessing, no other instruction can modify the memory location, including other Test-and-Set instructions
- Useful for implementing lock operations





Test-and-Set in RISC-V using lr/sc

- Example: RISC-V sequence for implementing a T&S at (s1)

Try:

```
li t2, 1
```

```
lr t1, s1
```

```
bne t1, x0, Try
```

```
sc t0, s1, t2
```

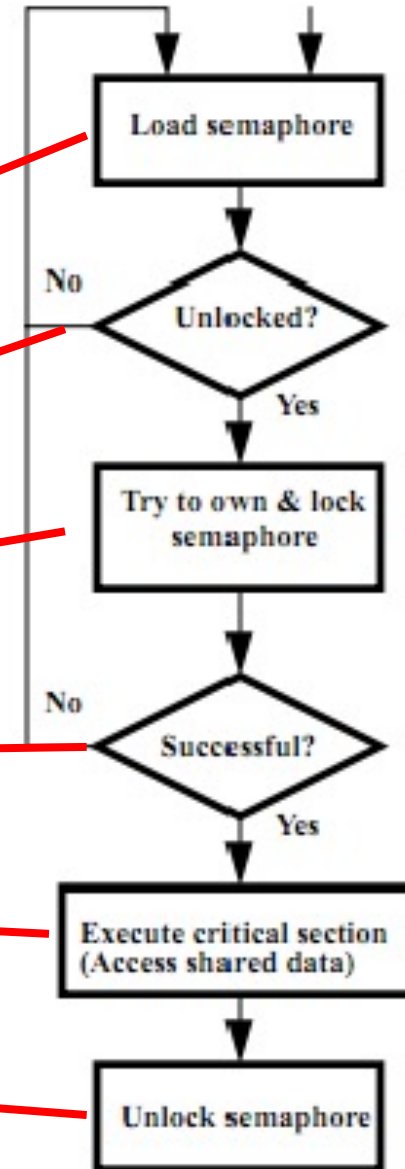
```
bne t0, x0, Try
```

Locked:

```
# critical section
```

Unlock:

```
sw x0, 0(s1)
```





Option 2: RISC-V Atomic Memory Operations (AMOs)

- Encoded with an R-type instruction format
 - swap, add, and, or, xor, max, min
 - **AMOSWAP** rd, rs2, (rs1)
 - **AMOADD** rd, rs2, (rs1)
- Take the value **pointed to** by rs1
 - Load it into rd aq(acquire) and rl(release) to insure *in order* execution
 - Apply the operation to that value with the contents in rs2
 - If rs2==rd, use the old value in rd
 - Store the result back to where rs1 is pointed to
- This allows atomic swap as a primitive
 - It also allows “reduction operations” that are common to be efficiently implemented



RISC-V Critical Section

- Assume that the lock is in memory location stored in register a0
- The lock is “set” if it is 1; it is “free” if it is 0 (it’s initial value)

```
li          t0, 1          # Get 1 to set lock
Try: amoswap.w.aq t1, t0, (a0) # t1 gets old lock value
                                # while we set it to 1
bnez        t1, Try        # if it was already 1, another
                                # thread has the lock,
                                # so we need to try again
... critical section goes here ...
amoswap.w.rl x0, x0, (a0) # store 0 in lock to release
```



Lock Synchronization

Broken Synchronization

```
while (lock != 0) ;
```

```
lock = 1;
```

```
// critical section
```

```
lock = 0;
```

Fix (lock is at location (a0))

```
li t0, 1
```

```
Try: amoswap.w.aq t1, t0, (a0)
```

```
bnez t1, Try
```

```
Locked:
```

```
# critical section
```

```
Unlock:
```

```
amoswap.w.rl x0, x0, (a0)
```




How to use


- Don't implement yourself!
- Use according library – e.g.:
 - pthread
 - C++:
 - `std::thread` C++11
 - `std::jthread` C++20
 - `std::mutex`; `std::lock_guard`; `std::scoped_lock`; `std::shared_lock`
 - `std::condition_variable`; `std::counting_semaphore`; `std::latch`; `std::barrier`
 - `std::promise`; `std::future`
 - Qt QThread
 - OpenMP

<https://en.cppreference.com/w/cpp/thread>



OpenMP Programming Model - Review

- **Fork - Join Model:**



- OpenMP programs begin as single process (*master thread*) and executes sequentially until the first parallel region construct is encountered
 - **FORK:** Master thread then creates a team of parallel threads
 - Statements in program that are enclosed by the parallel region construct are executed in parallel among the various threads
 - **JOIN:** When the team of threads complete the statements in the parallel region construct, they synchronize and terminate, leaving only the master thread



parallel Pragma and Scope - Review

- Basic OpenMP construct for parallelization:

```
#pragma omp parallel  
{  
    /* code goes here */  
}
```


- *Each* thread runs a copy of code within the block
- Thread scheduling is *non-deterministic*
- OpenMP default is *shared* variables
 - To make private, need to declare with pragma:

```
#pragma omp parallel private (x)
```




OpenMP Directives (Work-Sharing)


- These are defined *within* a **parallel** section



Shares iterations of a loop across the threads



Each section is executed by a separate thread



Serializes the execution of a thread



parallel Statement Shorthand

```
#pragma omp parallel
{
    #pragma omp for
    for(i=0; i<len; i++) { ... }
}
```

This is the only
directive in the
parallel section

can be shortened to:

```
#pragma omp parallel for
    for(i=0; i<len; i++) { ... }
```

- Also works for sections



Building Block: `for` loop

```
for (i=0; i<max; i++) zero[i] = 0;
```

- Breaks *for loop* into chunks, and allocate each to a separate thread
 - e.g. if `max = 100` with 2 threads:
assign 0-49 to thread 0, and 50-99 to thread 1
- Must have relatively simple “shape” for an OpenMP-aware compiler to be able to parallelize it
 - Necessary for the run-time system to be able to determine how many of the loop iterations to assign to each thread
- No premature exits from the loop allowed ← In general, don't jump outside of any pragma block
 - i.e. No `break`, `return`, `exit`, `goto` statements




Parallel `for` *pragma*

```
#pragma omp parallel for
```

```
for (i=0; i<max; i++) zero[i] = 0;
```

- Master thread creates additional threads, each with a separate execution context
- All variables declared outside for loop are shared by default, except for loop index which is *private* per thread (Why?)
- Implicit “barrier” synchronization at end of for loop
- Divide index regions sequentially per thread
 - Thread 0 gets 0, 1, ..., (max/n)-1;
 - Thread 1 gets max/n, (max/n)+1, ..., 2*(max/n)-1





OpenMP Example

```
1  /* clang -Xpreprocessor -fopenmp -lomp -o for for.c */
2
3  #include <stdio.h>
4  #include <omp.h>
5  int main()
6  {
7      omp_set_num_threads(4);
8      int a[] = { 0, 1, 2, 3, 4, 5, 6, 7, 8, 9 };
9      int N = sizeof(a)/sizeof(int);
10
11     #pragma omp parallel for
12     for (int i=0; i<N; i++) {
13         printf("thread %d, i = %2d\n",
14             omp_get_thread_num(), i);
15         a[i] = a[i] + 10 * omp_get_thread_num();
16     }
17
18     for (int i=0; i<N; i++) printf("%02d ", a[i]);
19     printf("\n");
20 }
```

```
$ gcc-5 -fopenmp for.c; ./a.out

% clang -Xpreprocessor -fopenmp -lomp -o for for.c;
./for

thread 0, i = 0
thread 1, i = 3
thread 2, i = 6
thread 3, i = 8
thread 0, i = 1
thread 1, i = 4
thread 2, i = 7
thread 3, i = 9
thread 0, i = 2
thread 1, i = 5
00 01 02 13 14 15 26 27 38 39
```

The call to find the maximum number of threads that are available to do work is `omp_get_max_threads()` (from `omp.h`).



OpenMP Timing

- Elapsed wall clock time:

```
double omp_get_wtime(void);
```


- Returns elapsed wall clock time in seconds
- Time is measured per thread, no guarantee can be made that two distinct threads measure the same time
- Time is measured from “some time in the past,” so subtract results of two calls to `omp_get_wtime` to get elapsed time



Matrix Multiply in OpenMP

```
// C[M][N] = A[M][P] * B[P][N]
start_time = omp_get_wtime();
#pragma omp parallel for private(tmp, j, k)
    for (i=0; i<M; i++){
        for (j=0; j<N; j++){
            tmp = 0.0;
            for( k=0; k<P; k++){
                /* C(i,j) = sum(over k) A(i,k) * B(k,j) */
                tmp += A[i][k] * B[k][j];
            }
            C[i][j] = tmp;
        }
    }
run_time = omp_get_wtime() - start_time;
```

Outer loop spread across n threads;
inner loops inside a single thread





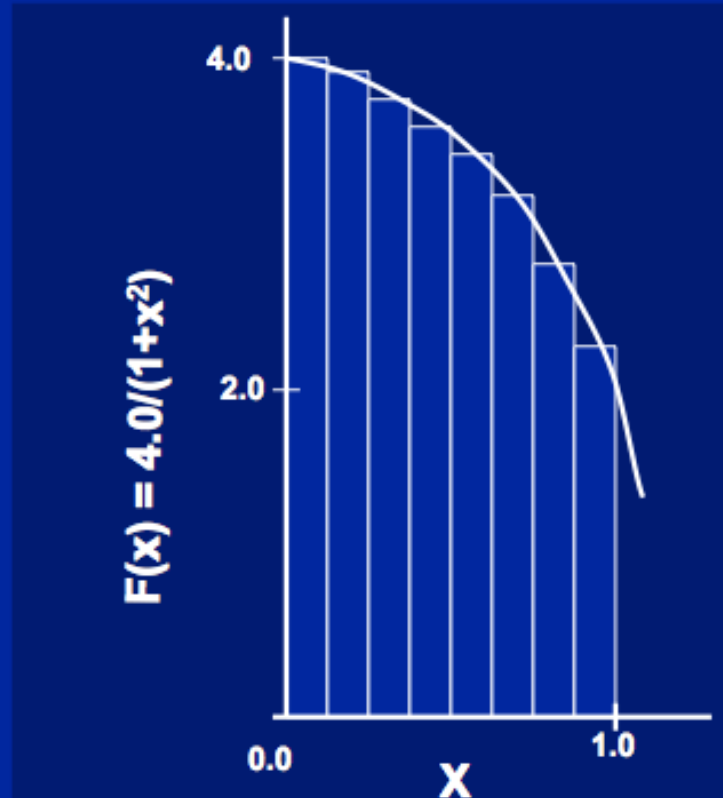
Notes on Matrix Multiply Example

- More performance optimizations available:
 - Higher *compiler optimization* (-O2, -O3) to reduce number of instructions executed
 - *Cache blocking* to improve memory performance
 - Using SIMD SSE instructions to raise floating point computation rate (*DLP*)



Example: Calculating π

Numerical Integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .



Sequential π

```
#include <stdio.h>

void main () {
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum = 0.0;
    for (int i=0; i<num_steps; i++) {
        double x = (i+0.5) *step;
        sum += 4.0*step/(1.0+x*x);
    }
    printf ("pi = %6.12f\n", sum);
}

pi = 3.142425985001
```

- Resembles π , but not very accurate
- Let's increase **num_steps** and parallelize



Parallelize (1) ...

```
#include <stdio.h>
```

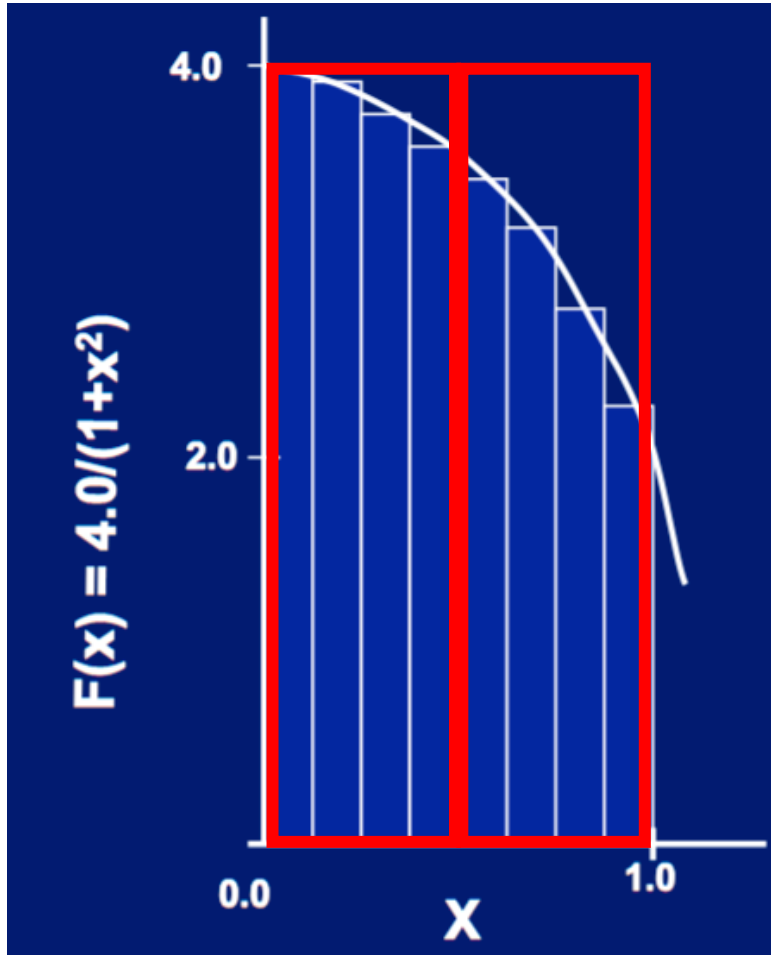
```
#include <omp.h>
```

```
void main () {  
    const long num_steps = 10;  
    double step = 1.0/((double)num_steps);  
    double sum = 0.0;  
    #pragma parallel for  
    for (int i=0; i<num_steps; i++) {  
        double x = (i+0.5) *step;  
        sum += 4.0*step/(1.0+x*x);  
    }  
    printf ("pi = %6.12f\n", sum);  
}
```

- Problem: each thread needs access to the shared variable **sum**
- Code runs sequentially ...



Parallelize (2) ...



1. Compute
 $\text{sum}[0]$ and $\text{sum}[1]$
in parallel

2. Compute
 $\text{sum} = \text{sum}[0] + \text{sum}[1]$
sequentially



Parallel π —Trial Run

```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 10;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        for (int i=id; i<num_steps; i+=NUM_THREADS) {
            double x = (i+0.5) * step;
            sum[id] += 4.0*step/(1.0+x*x);
            printf("i =%3d, id =%3d\n", i, id);
        }
    }
    double pi = 0;
    for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
    printf ("pi = %6.12f\n", pi);
}
```

```
i = 1, id = 1
i = 0, id = 0
i = 2, id = 2
i = 3, id = 3
i = 5, id = 1
i = 4, id = 0
i = 6, id = 2
i = 7, id = 3
i = 9, id = 1
i = 8, id = 0
pi = 3.142425985001
```

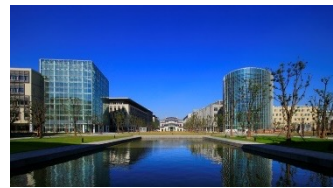
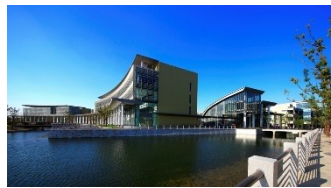



```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 4;
    const long num_steps = 1000000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        for (int i=id; i<num_steps; i+=NUM_THREADS) {
            double x = (i+0.5) *step;
            sum[id] += 4.0*step/(1.0+x*x);
            // printf("i =%3d, id =%3d\n", i, id);
        }
    }
    double pi = 0;
    for (int i=0; i<NUM_THREADS; i++) pi += sum[i];
    printf ("pi = %6.12f\n", pi);
}
```

pi =
3.141592653590

You verify how many
digits are correct ...



Can We Parallelize Computing sum?

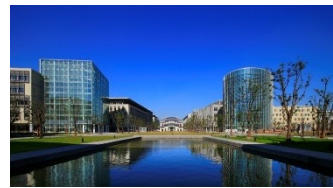
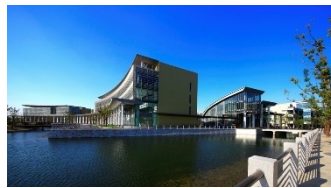
```
#include <stdio.h>
#include <omp.h>

void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        for (int i=id; i<num_steps; i+=NUM_THREADS) {
            double x = (i+0.5) *step;
            sum[id] += 4.0*step/(1.0+x*x);
        }
        pi += sum[id];
    }
    printf ("pi = %6.12f\n", pi);
}
```

Always looking for ways
to beat Amdahl's Law ...

Summation inside parallel section

- Insignificant speedup in this example, but ...
- **pi = 3.138450662641**
- Wrong! And value changes between runs?!
- What's going on?



What's Going On?

```
#include <stdio.h>
#include <omp.h>
```

```
void main () {
    const int NUM_THREADS = 1000;
    const long num_steps = 100000;
    double step = 1.0/((double)num_steps);
    double sum[NUM_THREADS];
    for (int i=0; i<NUM_THREADS; i++) sum[i] = 0;
    double pi = 0;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        for (int i=id; i<num_steps; i+=NUM_THREADS) {
            double x = (i+0.5) *step;
            sum[id] += 4.0*step/(1.0+x*x);
        }
        pi += sum[id];
    }
    printf ("pi = %6.12f\n", pi);
}
```

Can you resolve such a problem?

- Operation is really
$$pi = pi + sum[id]$$
- What if >1 threads reads current (same) value of **pi**, computes the sum, stores the result back to **pi**?
- Each processor reads same intermediate value of **pi**!
- Result depends on who gets there when
 - A “race” → result is not deterministic


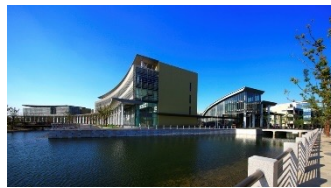


OpenMP Reduction

```
double avg, sum=0.0, A[MAX]; int i;  
#pragma omp parallel for private ( sum )  
for (i = 0; i <= MAX ; i++)  
    sum += A[i];  
avg = sum/MAX;  // bug
```

- *Problem is that we really want sum over all threads!*
- **Reduction**: specifies that, 1 or more variables that are private to each thread, are subject of reduction operation at end of parallel region:
reduction(operation:var) where
 - **Operation**: operator to perform on the variables (var) at the end of the parallel region : +, *, -, &, ^, |, &&, or ||.
 - **Var**: One or more variables on which to perform scalar reduction.

```
double avg, sum=0.0, A[MAX]; int i;  
#pragma omp for reduction(+ : sum)  
for (i = 0; i <= MAX ; i++)  
    sum += A[i];  
avg = sum/MAX;
```

parallel for, reduction

```
#include <omp.h>

#include <stdio.h>

static long num_steps = 100000;

double step;

void main () {

    int i;          double x, pi, sum = 0.0;

    step = 1.0 / (double)num_steps;

    #pragma omp parallel for private(x) reduction(+:sum)

    for (i=1; i<= num_steps; i++) {

        x = (i - 0.5) * step;

        sum = sum + 4.0 / (1.0+x*x);

    }

    pi = sum * step;

    printf ("pi = %6.12f\n", pi);


}
```

```
wangc@HP:~/TT$ gcc pi.c -o p -fopenmp
wangc@HP:~/TT$ ./p
pi = 3.141592653598
```




More on OpenMP


- These are defined *within* a `parallel` section



Shares iterations of a loop across the threads



Each section is executed by a separate thread



Serializes the execution of a thread

There are more, like `critical`, `barrier`, `atomic`, `master`, ... Try them by yourself.



上海
Shanghai

More

- Th

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char **argv) {
    int i = 0;
    omp_set_num_threads(4); // Maximum 4 threads
    #pragma omp parallel private(i)
    {
        printf("thread %d start\n", omp_get_thread_num());


        #pragma omp single
        {
            for (i = 0; i < 6; i++)
            {
                printf("Single, thread %d execute i = %d\n",
                    omp_get_thread_num(), i);
            }
        }
    }
}
```

single: code block executed by one thread only;
Other threads will wait;
Useful for thread-unsafe code;
Useful for I/O operations.

Each section is executed by a separate thread

Serializes the execution of a thread

, barrier, atomic, master, ... Try them by yourself.





上海
Shanghai

More

- Th

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char **argv) {
    int i = 0;
    omp_set_num_threads(4); // Maximum 4 threads
    #pragma omp parallel private(i)
    {
        printf("thread %d start\n", omp_get_thread_num());

        #pragma omp single
        {
            for (i = 0; i < 6; i++)
            {
                printf("Single, thread %d execute i = %d\n",
                    omp_get_thread_num(), i);
            }
        }
    }
}
```

single: code block executed by one thread only;
Other threads will wait;
Useful for thread-unsafe code;
Useful for I/O operations.

```
wangc@HP:~/TT$ gcc single.c -o s -fopenmp
wangc@HP:~/TT$ ./s
thread 3 start
Single, thread 3 execute i = 0
Single, thread 3 execute i = 1
Single, thread 3 execute i = 2
Single, thread 3 execute i = 3
Single, thread 3 execute i = 4
Single, thread 3 execute i = 5
thread 1 start
thread 2 start
thread 0 start
wangc@HP:~/TT$
```

master thread

RK

single

team





上
Sha

Mo

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char **argv) {
    int i = 0;
    omp_set_num_threads(4); // Maximum 4 threads
    #pragma omp parallel private(i)
    {
        printf("thread %d start\n", omp_get_thread_num());

        #pragma omp master
        {
            for (i = 0; i < 6; i++)
            {
                printf("Master, thread %d execute i = %d\n",
                    omp_get_thread_num(), i);
            }
            printf("Outside master, thread %d execute i =
%d\n",
                omp_get_thread_num(), i);
        }
    }
}
```

master Directive ensures that only the master threads executes instructions in the block. There is no implicit barrier, so other threads will not wait for master to finish


a separate thread

of a thread

barrier, atomic, master, ... Try them by yourself.



section





上
Sha

Mo

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char **argv) {
    int i = 0;
    omp_set_num_threads(4); // Maximum 4 threads
    #pragma omp parallel private(i)
    {
        printf("thread %d start\n", omp_get_thread_num());

        #pragma omp master
        {
            for (i = 0; i < 6; i++)
            {
                printf("Master, thread %d execute i = %d\n",
                    omp_get_thread_num(), i);
            }
            printf("Outside master, thread %d execute i = %d\n",
                omp_get_thread_num(), i);
        }
    }
}
```

master Directive ensures that only the master threads executes instructions in the block. There is no implicit barrier, so other threads will not wait for master to finish

section

↓ master thread

```
wangc@HP:~/TT$ gcc master.c -o m -fopenmp
wangc@HP:~/TT$ ./m
thread 2 start
Outside master, thread 2 execute i = 0
thread 1 start
Outside master, thread 1 execute i = 0
thread 3 start
Outside master, thread 3 execute i = 0
thread 0 start
Master, thread 0 execute i = 0
Master, thread 0 execute i = 1
Master, thread 0 execute i = 2
Master, thread 0 execute i = 3
Master, thread 0 execute i = 4
Master, thread 0 execute i = 5
Outside master, thread 0 execute i = 6
wangc@HP:~/TT$
```



Shanghai Tech University



Master

```
wangc@HP-Z2-G4:~/Works/TT$ ./p
thread 0 start
Master, thread 0 execute i = 0
Master, thread 0 execute i = 1
thread 2 start
Outside master, thread 2 execute i = 0
Master, thread 0 execute i = 2
Master, thread 0 execute i = 3
Master, thread 0 execute i = 4
Master, thread 0 execute i = 5
Outside master, thread 0 execute i = 6
thread 3 start
Outside master, thread 3 execute i = 0
thread 1 start
Outside master, thread 1 execute i = 0
```

```
printf("Outside m
%d\n",
omp_get_thread_num(), i);
}
```

master Directive ensures that only the master threads executes instructions in the block. There is no implicit barrier, so other threads will not wait for master to finish

```
4 threads
omp_get_thread_num());
i++)
ster, thread %d execute i = %d\n",
```



```
wangc@HP:~/TT$ gcc master.c -o m -fopenmp
wangc@HP:~/TT$ ./m
thread 2 start
Outside master, thread 2 execute i = 0
thread 1 start
Outside master, thread 1 execute i = 0
thread 3 start
Outside master, thread 3 execute i = 0
thread 0 start
Master, thread 0 execute i = 0
Master, thread 0 execute i = 1
Master, thread 0 execute i = 2
Master, thread 0 execute i = 3
Master, thread 0 execute i = 4
Master, thread 0 execute i = 5
Outside master, thread 0 execute i = 6
wangc@HP:~/TT$
```



And in Conclusion, ...

- Multiprocessor/Multicore uses Shared Memory
 - Cache coherency implements shared memory even with multiple copies in multiple caches
 - False sharing a concern; watch block size!
 - To be covered with “Advanced caches” :-)
- OpenMP as simple parallel extension to C
 - Threads, Parallel for, private, reductions ...
 - \approx C: small so easy to learn, but not very high level and it's easy to get into trouble
 - Much we didn't cover – including other synchronization mechanisms (locks, etc.)