CS110 Computer Architecture
Dependability and RAID

Chundong Wang & Siting Liu

SIST, ShanghaiTech

TOAST LAB
吐 司 实 验 室

Review

- DMA
- Disk I/Os
- Flash memory
- Network
 - Abstraction
  - Protocols

Protocol Family Concept

- *Protocol*: packet structure and control commands to manage communication

- *Protocol families (suites)*: a set of cooperating protocols that implement the network stack

- Key to protocol families is that communication occurs logically at the same level of the protocol, called peer-to-peer…

  …but is implemented via services at the next lower level

- Encapsulation: carry higher level information within lower level "envelope"
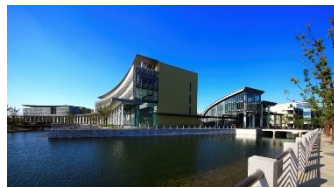
Inspiration…

Dear Jack,

It is ready.

--Pony

Inspiration…

- CEO A writes letter to CEO B
  - Folds letter and hands it to assistant
- Assistant:
  - Puts letter in envelope with CEO B's full name
  - Takes to China Post
- China Post Office
  - Puts letter in larger envelope
  - Puts name and street address on China Post envelope
  - Puts package on China Post delivery truck
- China Post delivers to other company

The Path of the Letter

"Peers" on each side understand the same things
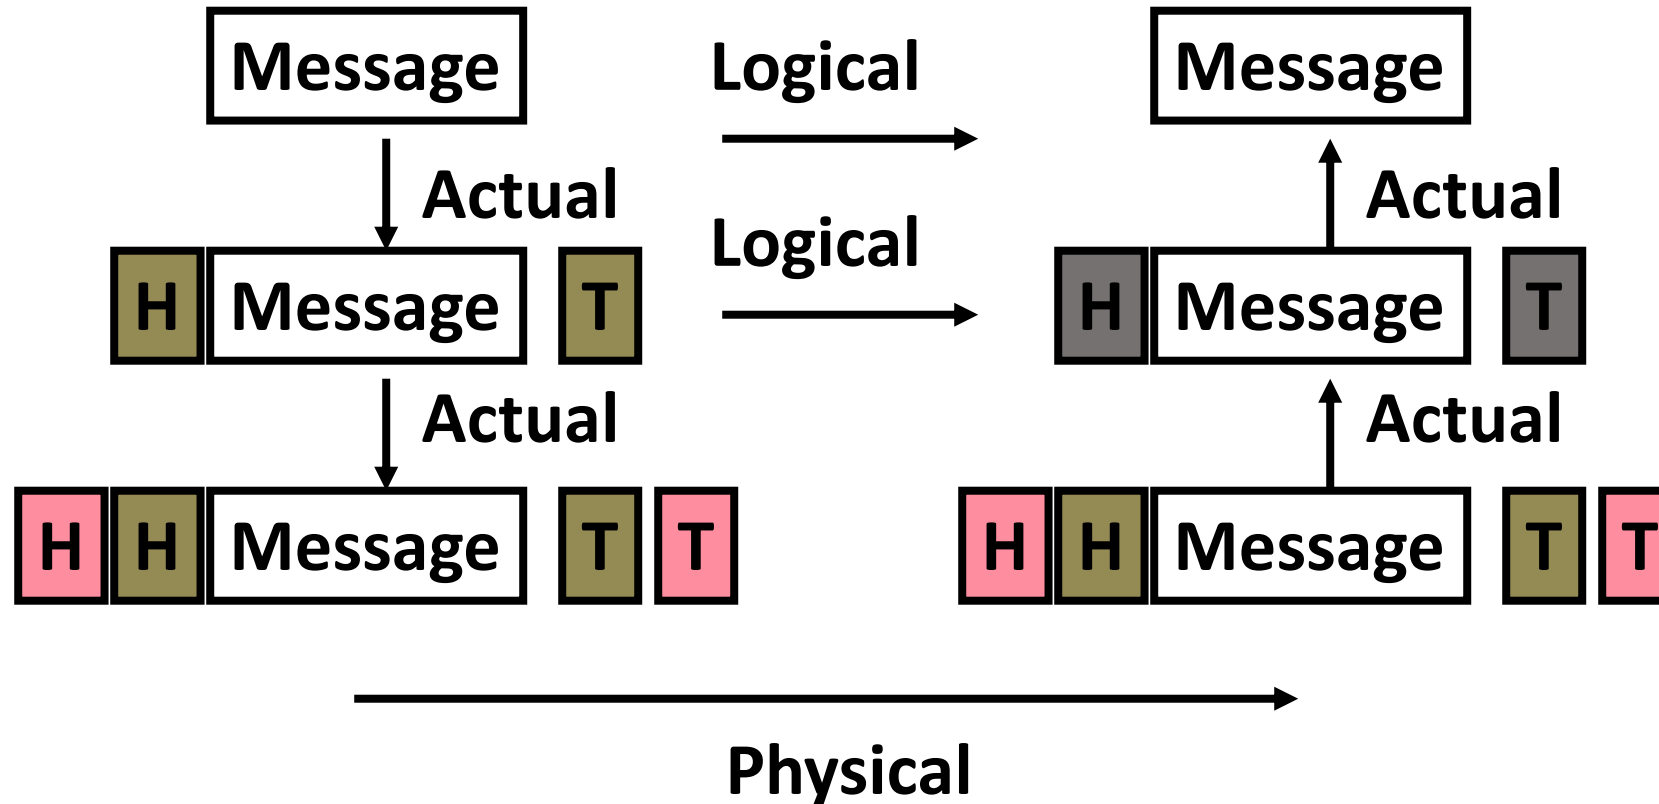No one else needs to
Lowest level has most packaging

CEO   Letter   CEO

Aide   Envelope   Aide

China Post ──────────── China Post ──────────► China Post

The Path of the Letter

"Peers" on each side understand the same things
No one else needs to
Lowest level has most packaging

| CEO	Semantic Content	CEO
Aide	Identity	Aide
China Post	Location	China Post
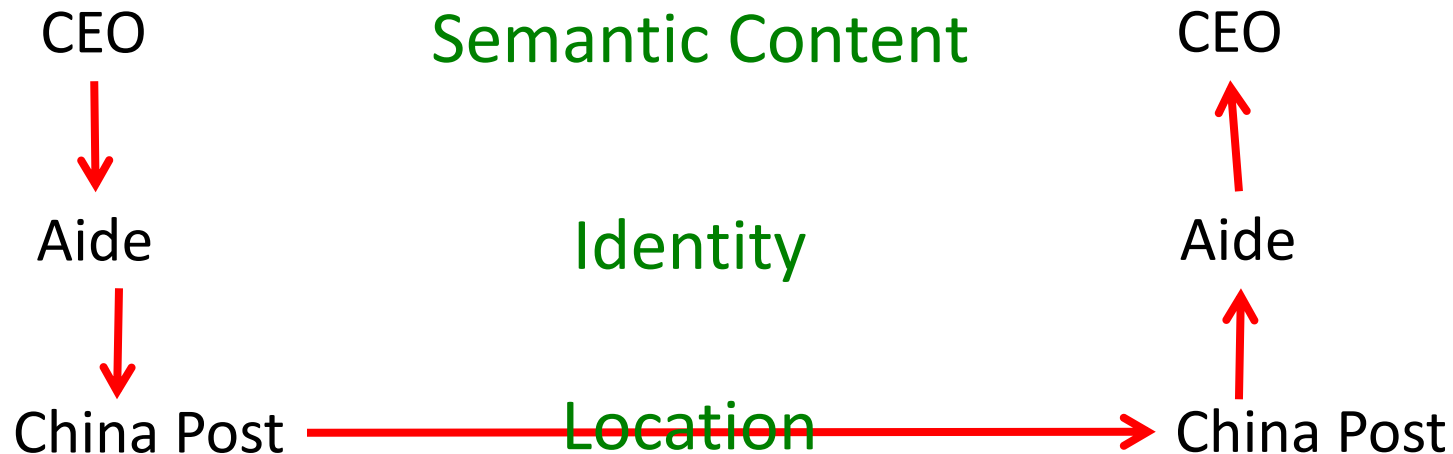
Protocol Family Concept

*Each lower level of stack "encapsulates" information from layer above by adding header and trailer.*

Most Popular Protocol for Network of Networks

- Transmission Control Protocol/Internet Protocol (TCP/IP)
- This protocol family is the basis of the Internet, a WAN (wide area network) protocol
  - IP makes best effort to deliver
    - Packets can be lost, corrupted
  - TCP guarantees delivery
  - TCP/IP so popular it is used even when communicating locally: even across homogeneous LAN (local area network)
  - UDP/IP: video or sound streaming; video call….
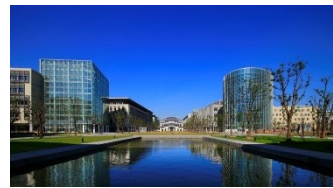
TCP/IP packet, Ethernet packet, protocols

- Application sends message

  - **TCP breaks into 64KiB segments, adds 20B header**

  - **IP adds 20B header, sends to network**

  - **If Ethernet, broken into 1500B packets with headers, trailers**

Ethernet Hdr

IP Header

TCP Header

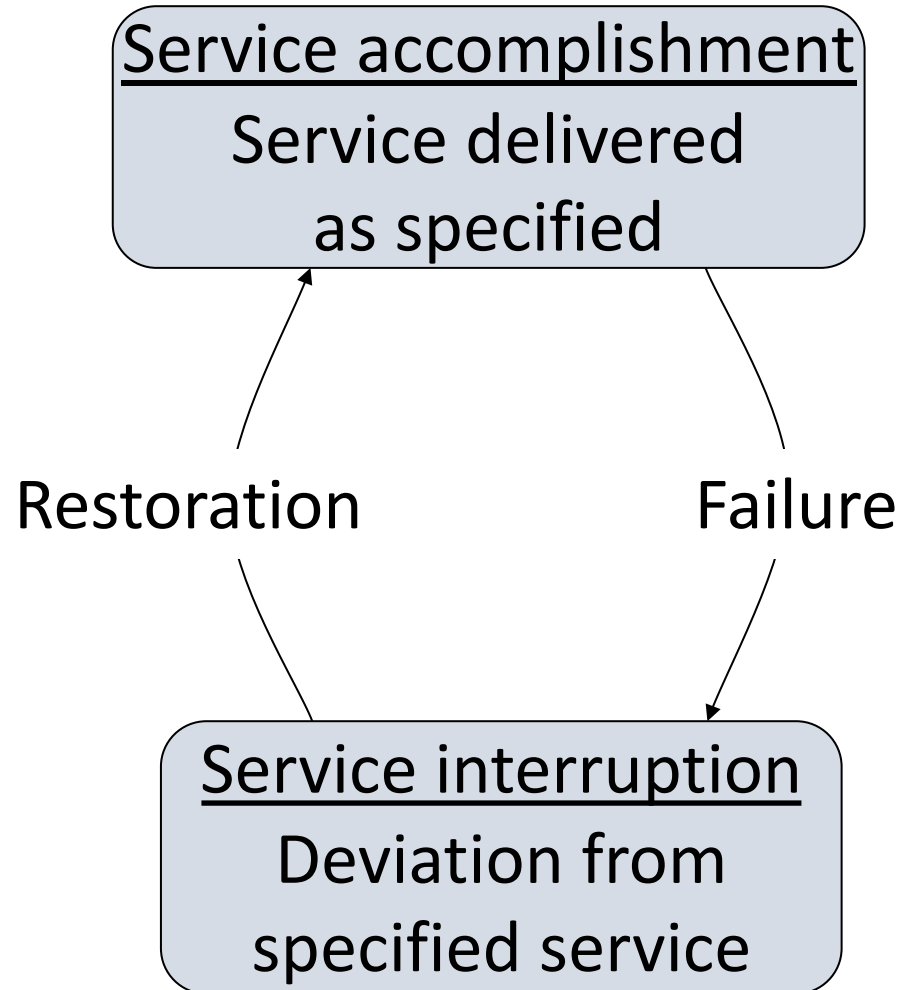**EHIP Data**

TCP data

Message

Ethernet Hdr
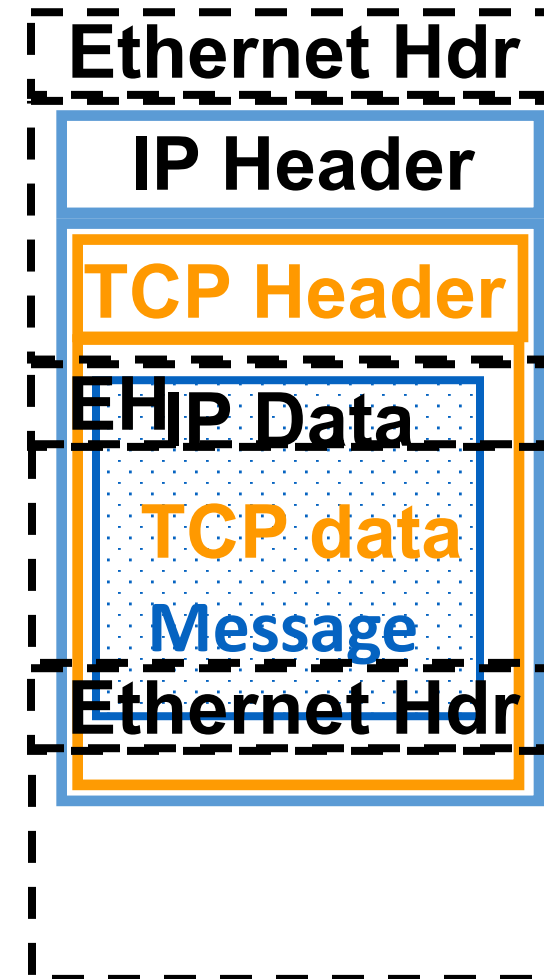
Great Idea #6: Dependability via Redundancy

- Applies to everything from datacenters to memory
  - Redundant datacenters so that can lose 1 datacenter but Internet service stays online
  - Redundant routes so can lose nodes but Internet doesn't fail
  - Redundant memory bits of so that can lose 1 bit but no data (Error Correcting Code/ECC Memory)
  - Redundant disks so that can lose 1 disk but not lose data (Redundant Arrays of Independent Disks/RAID)

Dependability

Service accomplishment
Service delivered
as specified

Restoration   Failure

Service interruption
Deviation from
specified service

- Fault: failure of a component
  - May or may not lead to system failure

Dependability via Redundancy:
Time vs. Space

- *Spatial Redundancy* – replicated data or check information or hardware to handle hard and soft (transient) failures

- *Temporal Redundancy* – redundancy in time (retry) to handle soft (transient) failures
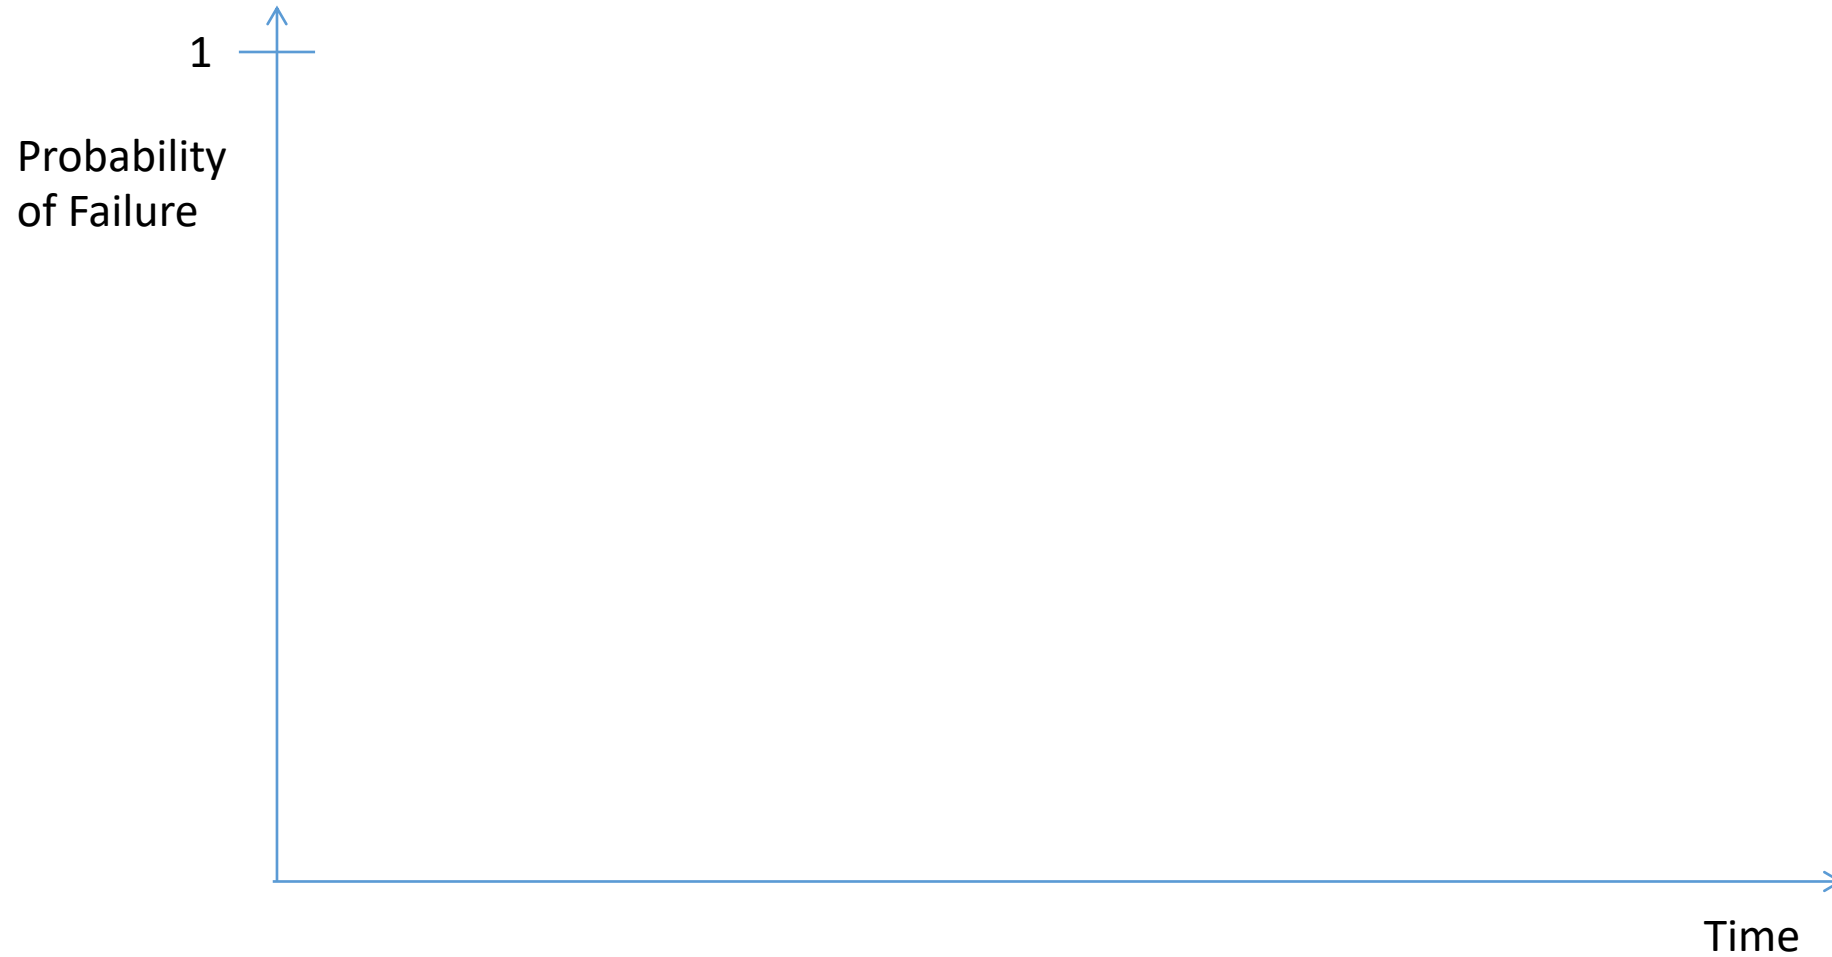
Dependability Measures

- Reliability: Mean Time To Failure (MTTF)

- Service interruption: Mean Time To Repair (MTTR)

- Mean time between failures (MTBF)
  - MTBF = MTTF + MTTR

- Availability = $\dfrac{MTTF}{MTTF+MTTR}$

- Improving Availability
  - Increase MTTF: More reliable hardware/software + Fault Tolerance
  - Reduce MTTR: improved tools and processes for diagnosis and repair
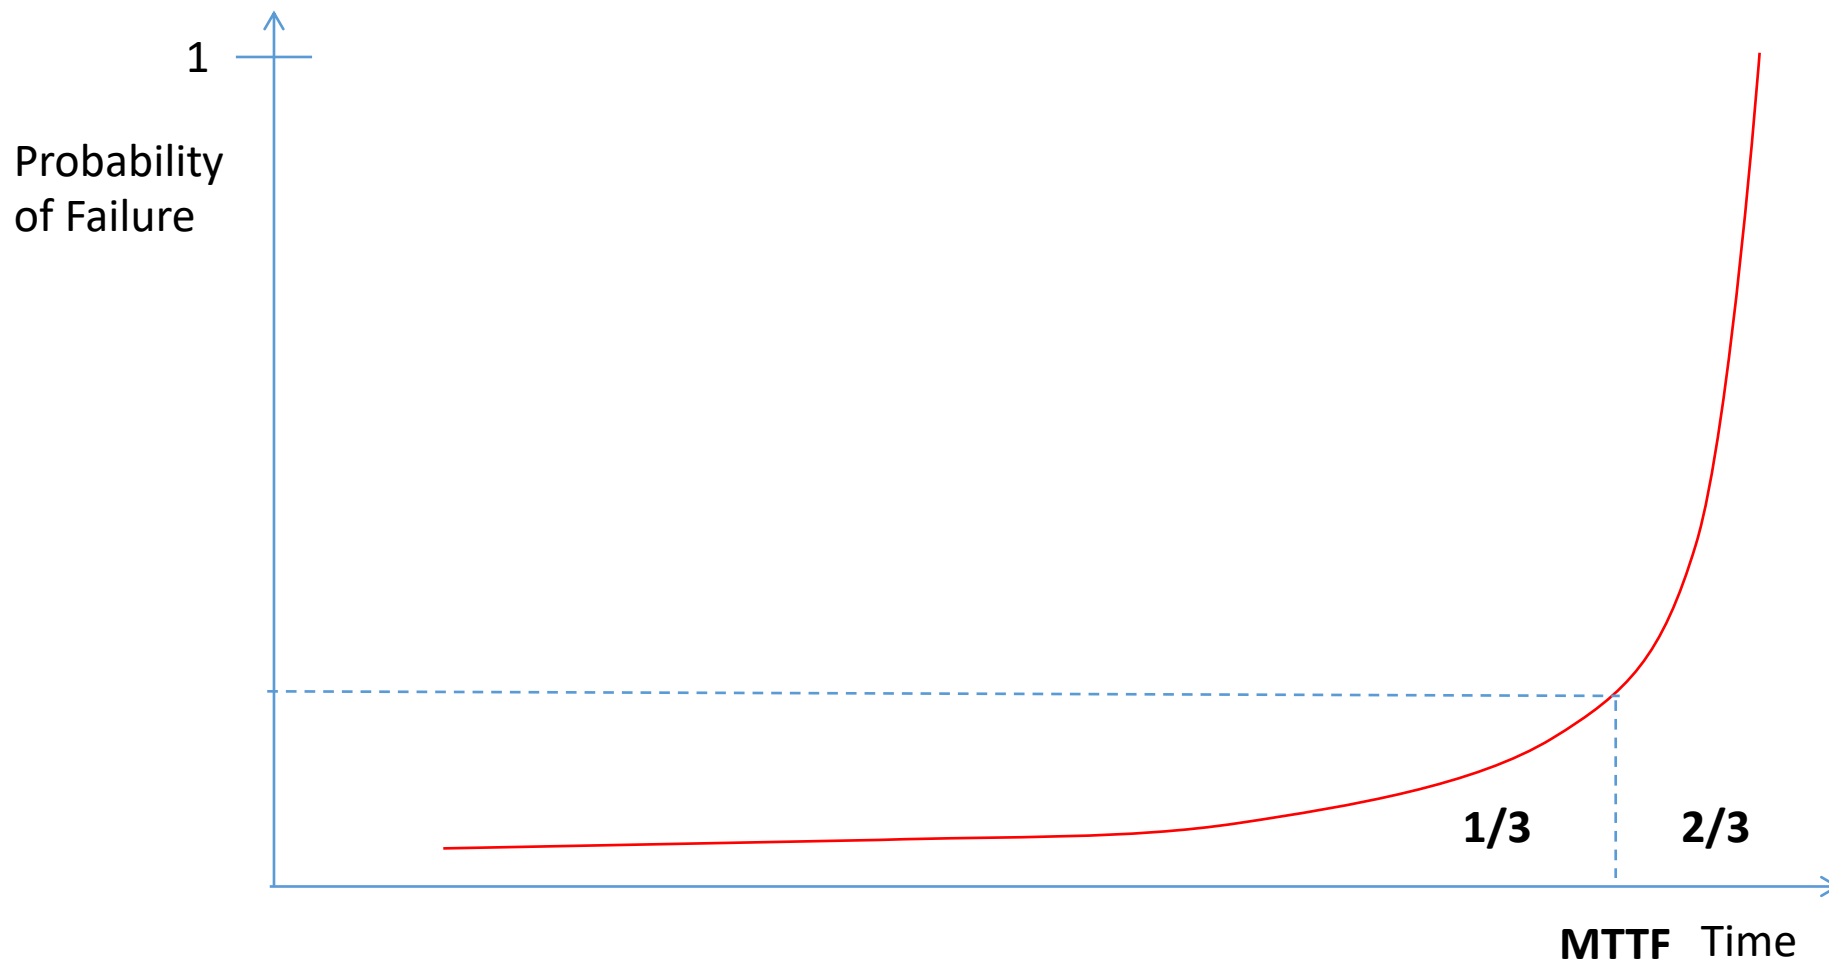
Understanding MTTF

Probability
of Failure

1

Time
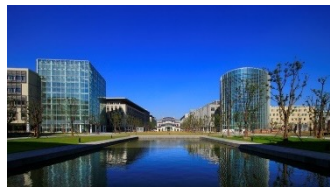
Understanding MTTF

Availability Measures

- Availability = $\dfrac{MTTF}{MTTF+MTTR}$ as %
  - MTTF, MTBF usually measured in hours
- Since hope rarely down, shorthand is "number of 9s of availability per year"
- 1 nine: 90% => 36 days of repair/year
- 2 nines: 99% => 3.6 days of repair/year
- 3 nines: 99.9% => 526 minutes of repair/year
- 4 nines: 99.99% => 53 minutes of repair/year
- 5 nines: 99.999% => 5 minutes of repair/year

Reliability Measures

- Another is average number of failures per year: Annualized Failure Rate (AFR)
  - E.g., 1000 disks with 100,000 hours MTTF
  - 365 days * 24 hours = 8760 hours
  - (1000 disks * 8760 hrs/year) / 100,000 = 87.6 failed disks per year on average
  - 87.6/1000 = 8.76% annual failure rate
- Google's 2007 study* found that actual AFRs for individual drives ranged from 1.7% for first year drives to over 8.6% for three-year old drives

*research.**google**.com/archive/disk_failures.pdf

Dependability Design Principle

- Design Principle: No single points of failure
  - "Chain is only as strong as its weakest link"
  - *Achilles' Heel*

- Dependability Corollary of Amdahl's Law
  - Doesn't matter how dependable you make one portion of system
  - Dependability limited by part you do not improve

Error Detection/Correction Codes

- Memory systems generate errors (accidentally flipped-bits)
  - DRAMs store very little charge per bit
  - "Soft" errors occur occasionally when cells are struck by alpha particles or other environmental upsets
  - "Hard" errors can occur when chips permanently fail
  - Problem gets worse as memories get denser and larger

- Memories protected against failures with EDC/ECC

- Extra bits are added to each data-word
  - Used to detect and/or correct faults in the memory system
  - Each data word value mapped to unique *code word*
  - A fault changes valid code word to invalid one, which can be detected

Block Code Principles

- Hamming distance = difference in # of bits

- p = 0<u>1</u>1<u>0</u>11, q = 0<u>0</u>1<u>1</u>11, Ham. distance (p,q) = 2

- p = 011011,
q = 110001,
distance (p,q) = ?

- Can think of extra bits as creating a code with the data
  - There is Ham. distance between codes

Richard Hamming, 1915-98
Turing Award Winner

25

Parity

- Parity bits are added to a word to make it
  - either odd: odd numbers of '1'
  - or even: even number of '1'

- Let us add one parity bit to three-bit word

Odd Parity		Even Parity	
000	0001	000	0000
100	1000	100	1001
101	1011	101	1010
111	1110	111	1111

Parity: Simple Error-Detection Coding

- Each data value, before it is written to memory is "tagged" with an extra bit to force the stored word to have *even parity*:

$$b_7 b_6 b_5 b_4 b_3 b_2 b_1 b_0$$

$\rightarrow p$

- Each word, as it is read from memory is "checked" by finding its parity (including the parity bit).

$$b_7 b_6 b_5 b_4 b_3 b_2 b_1 b_0 \quad p$$

c
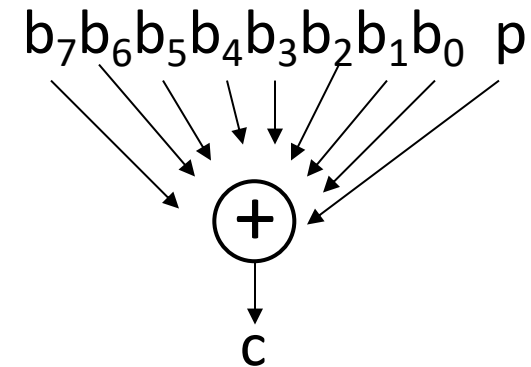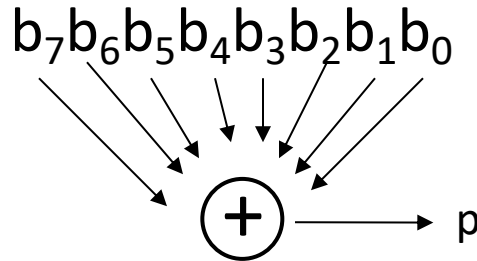
- A non-zero parity check indicates an error occurred:
  - 2 errors (on different bits) are not detected
  - nor any even number of errors, just odd numbers of errors are detected
- Minimum Hamming distance of valid parity codes is 2

Parity Example

- Data 0101 0101
- 4 ones, even parity now
- Write to memory:
  0101 0101 0
  to keep parity even
- Data 0101 0111
- 5 ones, odd parity now
- Write to memory:
  0101 0111 1
  to make parity even

- Read from memory
  0101 0101 0
- 4 ones => even parity, so no error
- Read from memory
  1101 0101 0
- 5 ones => odd parity,
  so error
- What if error in parity bit?

Suppose Want to Correct 1 Error?

- Richard Hamming came up with simple to understand mapping to allow Error Correction at minimum distance of 3
    - Single error correction, double error detection
- Called "Hamming ECC"
    - Worked weekends on relay computer with unreliable card reader, frustrated with manual restarting
    - Got interested in error correction; published 1950
    - R. W. Hamming, "Error Detecting and Correcting Codes," *The Bell System Technical Journal*, Vol. XXVI, No 2 (April 1950) pp 147-160.

Detecting/Correcting Code Concept

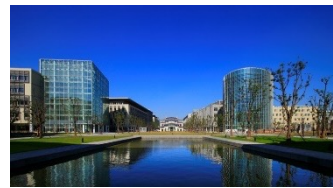Space of possible bit patterns (2^N)

Error changes bit pattern to non-code

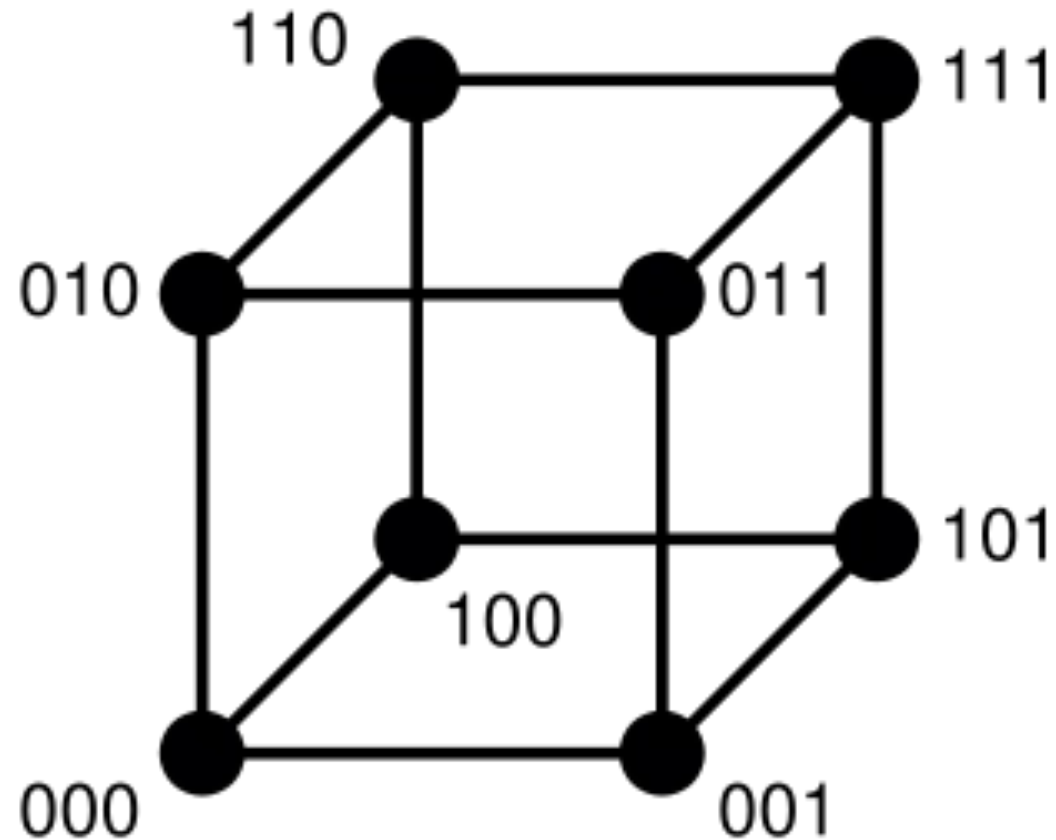Sparse population of valid code words ($2^M \ll 2^N$)
- with identifiable signature

- Detection: bit pattern fails codeword check

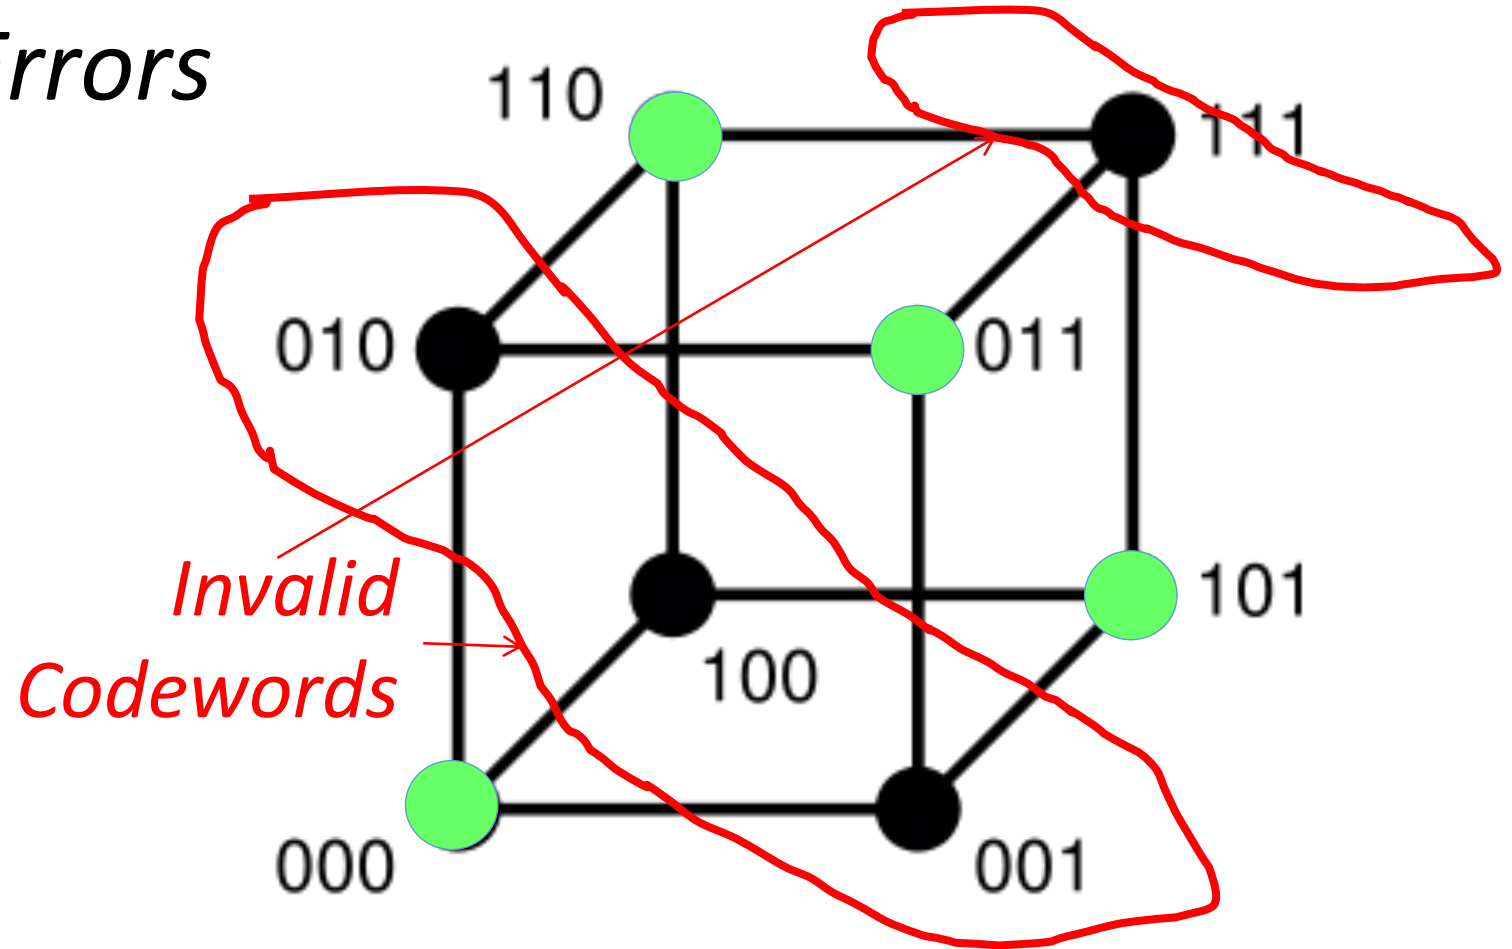- Correction: map to nearest valid code word

Hamming Distance: 8 code words

Hamming Distance 2: Detection
Detect Single Bit Errors

- No 1 bit error goes to another valid codeword
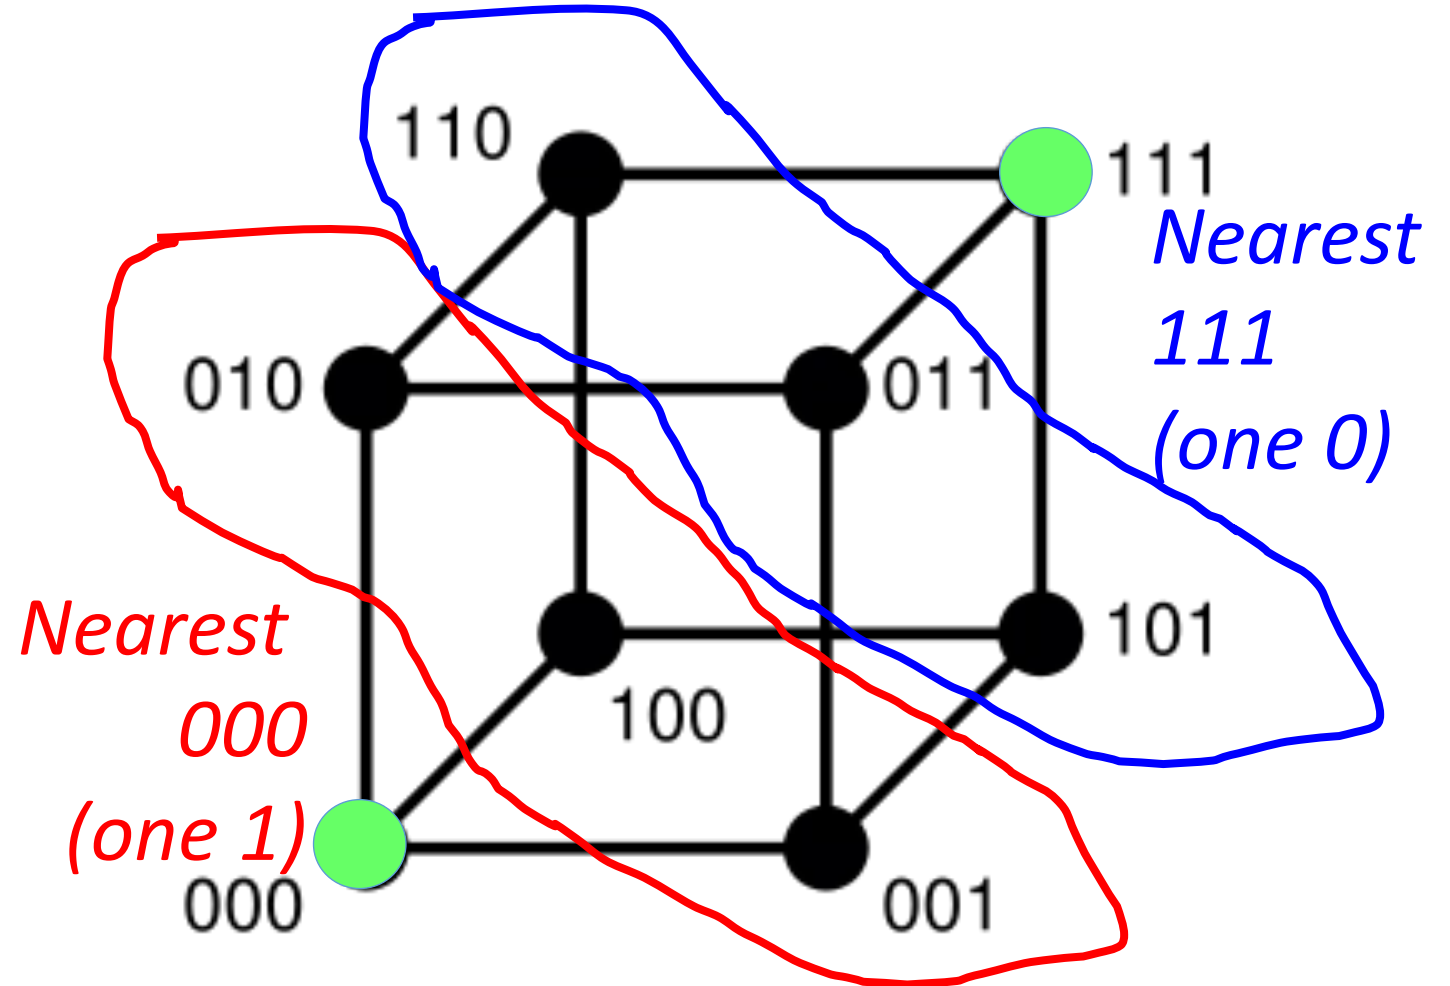- ½ codewords are valid

Invalid Codewords

Hamming Distance 3: Correction

Correct Single Bit Errors,
 Detect Double Bit Errors

- No 2 bit error goes to another valid codeword; 1 bit error near
- 1/4 codewords are valid

Nearest 111 (one 0)

Nearest 000 (one 1)

Hamming Error Correction Code

- Use of extra parity bits to allow the position identification of a single error

1. Mark all bit positions that are powers of 2 as parity bits (positions 1, 2, 4, 8, 16, …)
   - Start numbering bits at 1 at left (not at 0 on right)

2. All other bit positions are data bits
   (positions 3, 5, 6, 7, 9, 10, 11, 12, 13, 14, 15, …)

3. Each data bit is covered by 2 or more parity bits

Hamming ECC

4. The position of parity bit determines sequence of data bits that it checks

- Bit 1 (0001_2): checks bits (1,3,5,7,9,11,…)
  - Bits with <u>least</u> significant bit of address = 1
- Bit 2 (0010_2): checks bits (2,3,6,7,10,11,14,15,…)
  - Bits with <u>2nd</u> least significant bit of address = 1
- Bit 4 (0100_2): checks bits (4-7, 12-15, 20-23, …)
  - Bits with <u>3rd</u> least significant bit of address = 1
- Bit 8 (1000_2): checks bits (8-15, 24-31, 40-47 ,…)
  - Bits with <u>4th</u> least significant bit of address = 1

Graphic of Hamming Code

Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11
Parity bit coverage	p1	X		X		X		X		X		X		X		X
	p2		X	X			X	X			X	X			X	X
	p4				X	X	X	X					X	X	X	X
	p8								X	X	X	X	X	X	X	X

- http://en.wikipedia.org/wiki/Hamming_code

Hamming ECC

5. Set parity bits to create even parity for each group

- A byte of data: 10011010

- Create the coded word, leaving spaces for the parity bits:

- $\_\,\_\,1\,\_\,0\,0\,1\,\_\,1\,0\,1\,0$
  1 2 3 4 5 6 7 8 9 A B C

- Calculate the parity bits

Hamming ECC

_ _ 1 _ 0 0 1 _ 1 0 1 0

- Position 1 checks bits 1, 3, 5, 7, 9, 11:

  ? _ **1** _ **0** 0 **1** _ **1** 0 **1** 0. set position 1:

  0 _ **1** _ **0** 0 **1** _ **1** 0 **1** 0

- Position 2 checks bits 2, 3, 6, 7, 10, 11:

  0 **?** **1** _ 0 **0** **1** _ 1 **0** **1** 0. set position 2:

  0 **1** **1** _ 0 **0** **1** _ 1 **0** **1** 0

- Position 4 checks bits 4, 5, 6, 7, 12:

  0 1 1 **?** **0** **0** **1** _ 1 0 1 **0**. set position 4:

  0 1 1 **1** **0** **0** **1** _ 1 0 1 **0**

- Position 8 checks bits 8, 9, 10, 11, 12:

  - 0 1 1 1 0 0 1 **?** **1** **0** **1** **0**. set position 8:
  - 0 1 1 1 0 0 1 **0** **1** **0** **1** **0**
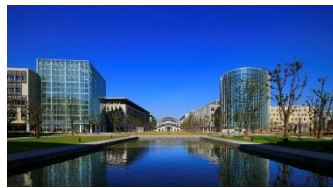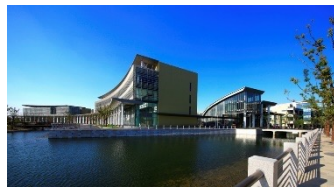
Hamming ECC

- Final code word: 0111001<u>0</u>1010
- Data word: 1   001 1010

Hamming ECC Error Check

- Suppose receive

  011100101110

  0 1 1 1 0 0 1 0 1 1 1 0

Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11
Parity bit coverage	p1	X		X		X		X		X		X		X		X
	p2		X	X			X	X			X	X			X	X
	p4				X	X	X	X					X	X	X	X
	p8								X	X	X	X	X	X	X	X

Hamming ECC Error Check

- Suppose receive

```
011100101110
0 1 0 1 1 1   √
 11   01   11    X-Parity 2 in error
    1001      0 √
          01110  X-Parity 8 in error
```

- *Implies position 8+2=10 is in error*

```
011100101110
```

Hamming ECC Error Correct

- Flip the incorrect bit …
  0111100101010

- Double check
  0111100101010
  0 1 0 1 1 1 √
   11 01   01 √
      1001   0 √
         01010 √

Hamming ECC Error Detect

- Suppose receive

```
01010000101010
0  0  0  0  1  1    √
 10   00   01    √
  1000      0  X
     01010  √
```

Two errors can be detected, but not correctable

How about ≥3 bits error?

Cyclic Redundancy Check

- Parity is not powerful enough to detect long runs of errors (also known as *burst errors*)

- Better Alternative: *Reed-Solomon Codes*
  - Used widely in CDs, DVDs, Magnetic Disks
  - RS(255,223) with 8-bit symbols: each codeword contains 255 code word bytes (223 bytes are dat
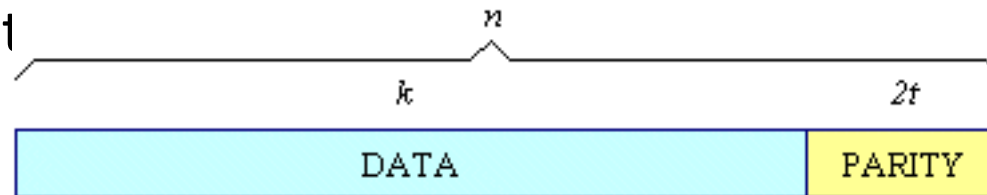
  - For this code: n = 255, k = 223, s = 8, 2t = 32, t = 16
  - Decoder can correct any errors in up to 16 bytes anywhere in the codeword

RAID: Redundancy for Disks

- Why we still worry about disks?
  - Trade-off: price, capacity, density, etc.
  - When you need storage space in petabytes (PB) or exabytes (EB)
    - 1 PB = 1024 TB
    - 1 EB = 1024 PB
  - Do not forget that flash-based SSDs also fail
    - Limited program/erase cycles ← wear levelling

RAID: Redundant Arrays of (Inexpensive) Disks

- Files are "striped" across multiple disks

- Redundancy yields high data availability
  - Availability: service still provided to user, even if some components failed

- Disks will still fail

- Contents reconstructed from data redundantly stored in the array
  - ➔ Capacity penalty to store redundant info
  - ➔ Bandwidth penalty to update redundant info

RAID 0: Striping

- RAID 0 provides **no** fault tolerance or redundancy
  - Striping, or disk spanning
  - High performance
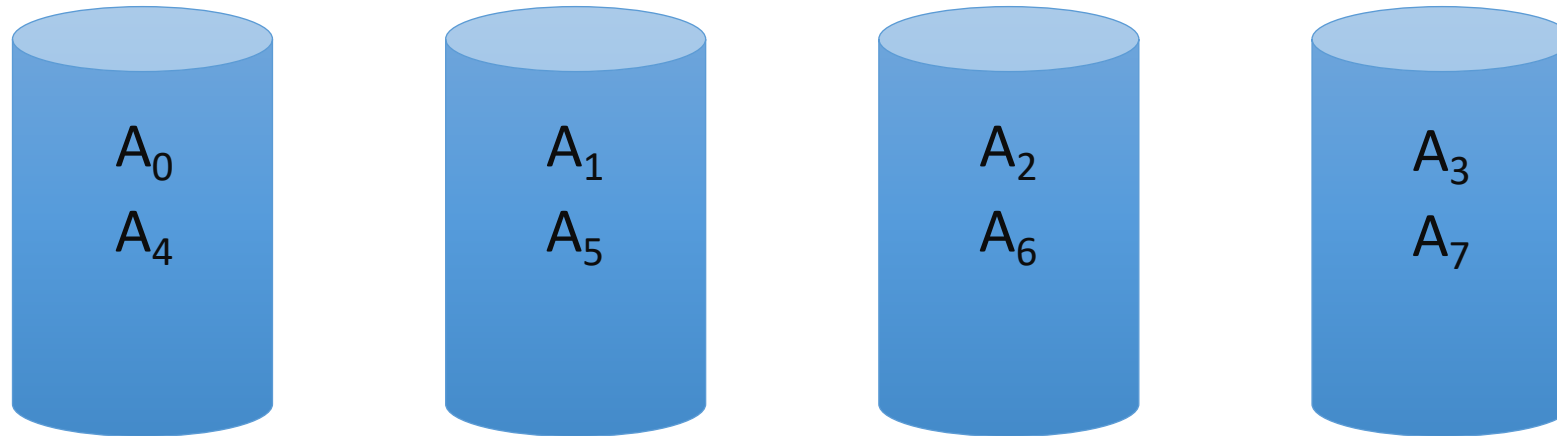
RAID 1: Disk Mirroring/Shadowing
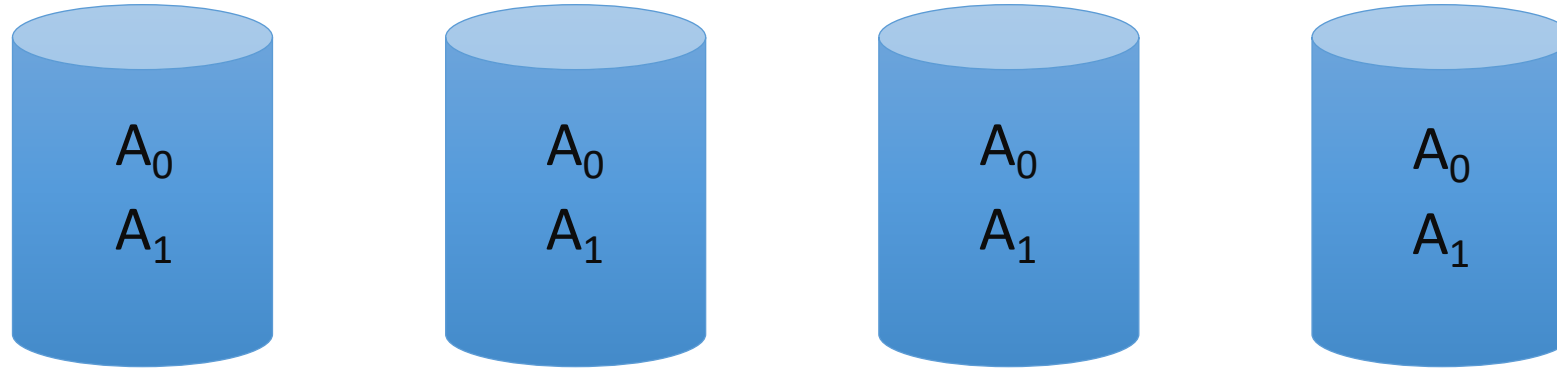
- Each disk is fully duplicated onto its "mirror(s)"
  - Very high availability can be achieved
- Bandwidth sacrifice on write:
  - Logical write = N physical writes
  - Reads may be optimized
- Most expensive solution: 100% capacity overhead
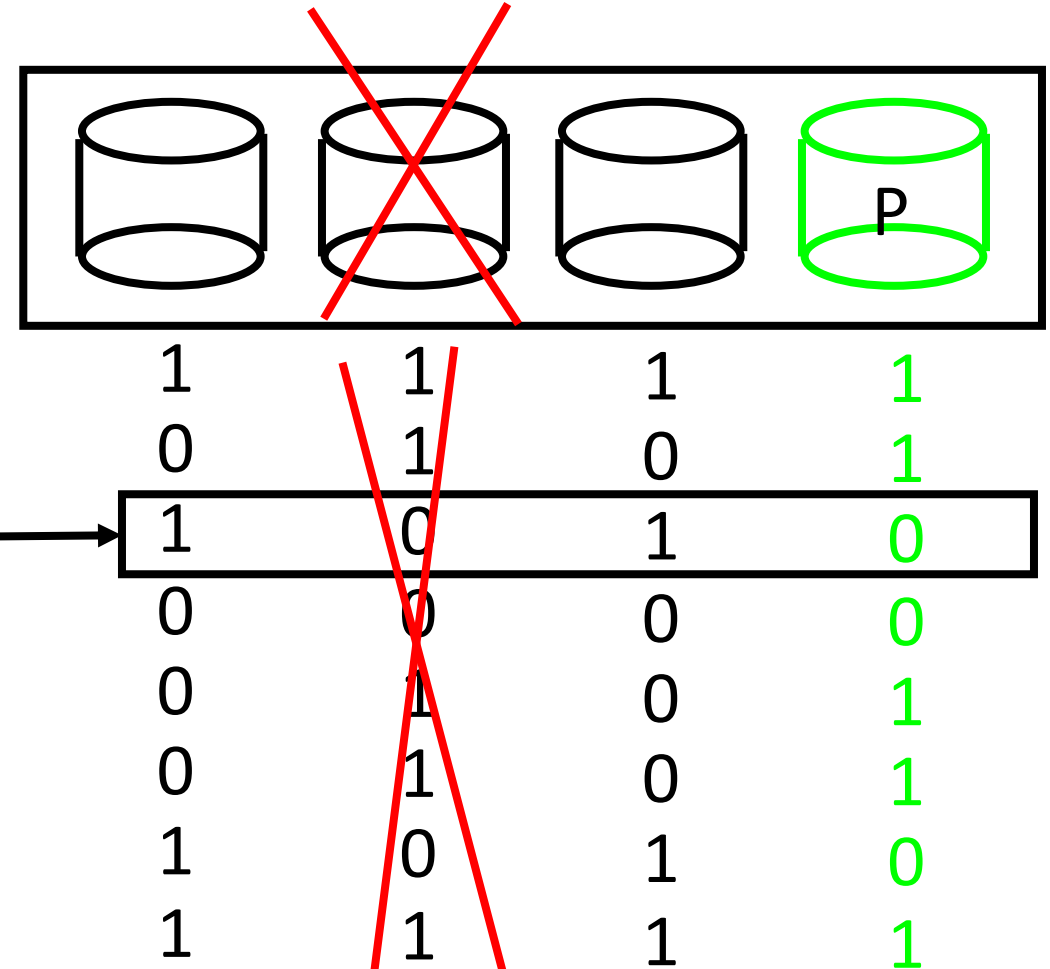- RAID 10 (striped mirrors), RAID 01 (mirrored stripes):
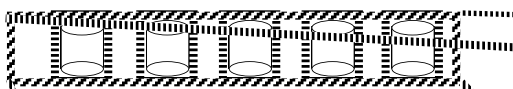 - Combinations of RAID 0 and 1.
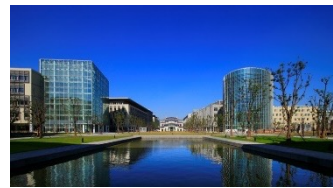
RAID 3: Parity Disk


```
10010011
11001101
10010011
. . .
```
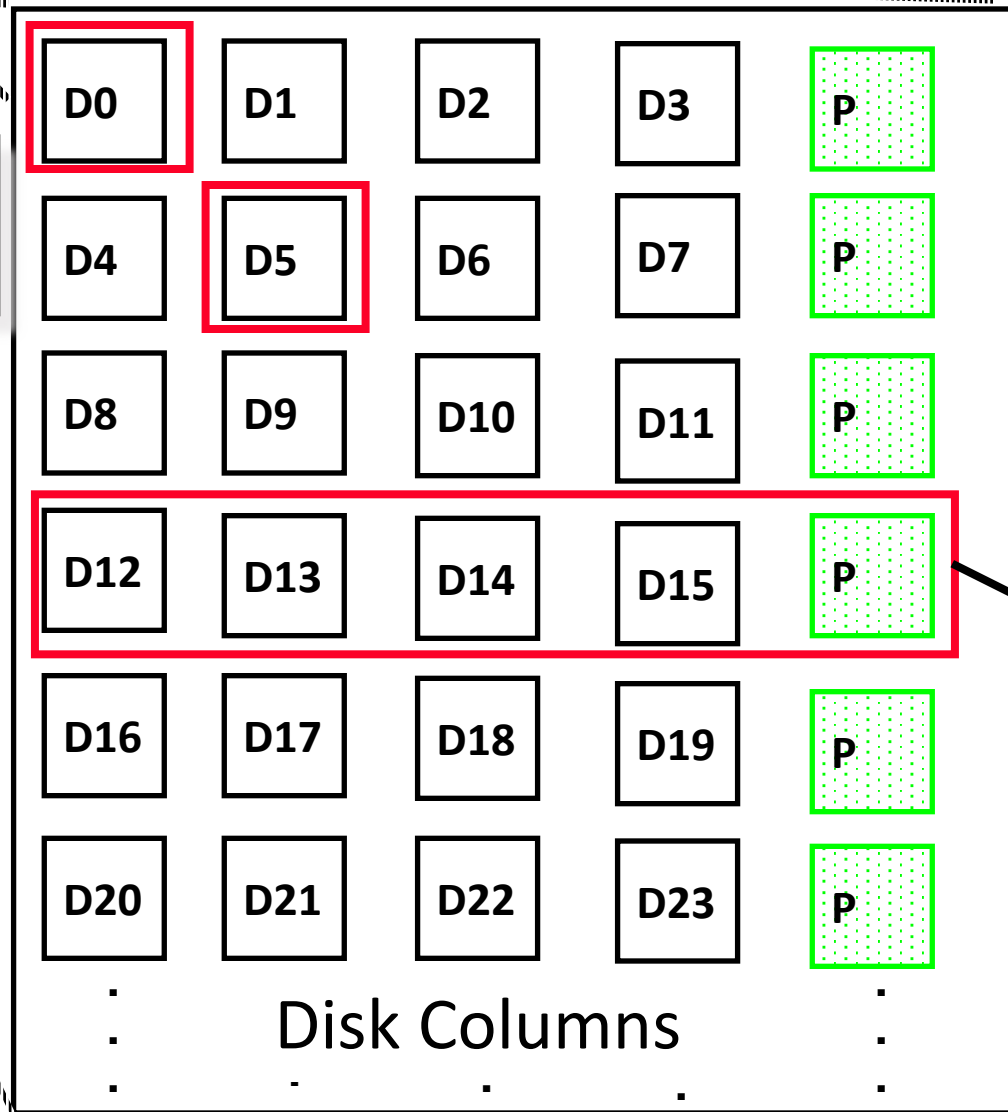logical record

Striped physical records

P contains sum of other disks per stripe mod 2 ("parity")
If disk fails, subtract P from sum of other disks to find missing information

Insides of 5 disks

Example: small read D0 & D5, large write D12-D15

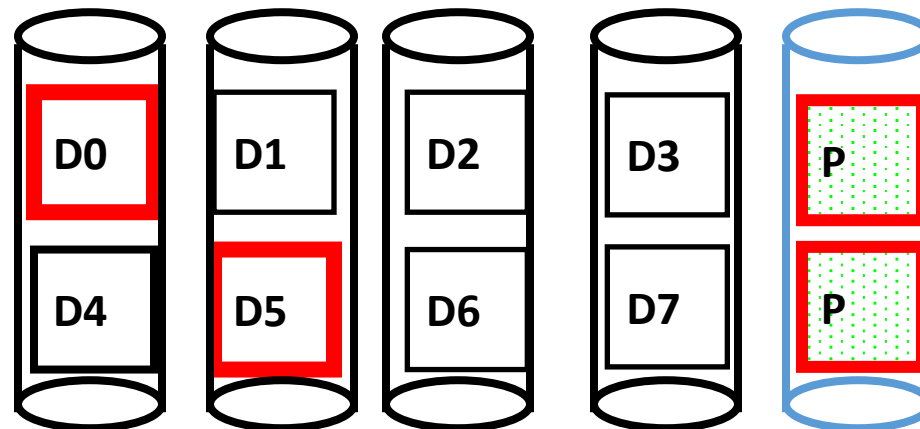| D0 | D1 | D2 | D3 | P |
D4	D5	D6	D7	P
D8	D9	D10	D11	P
D12	D13	D14	D15	P
D16	D17	D18	D19	P
D20	D21	D22	D23	P

Disk Columns

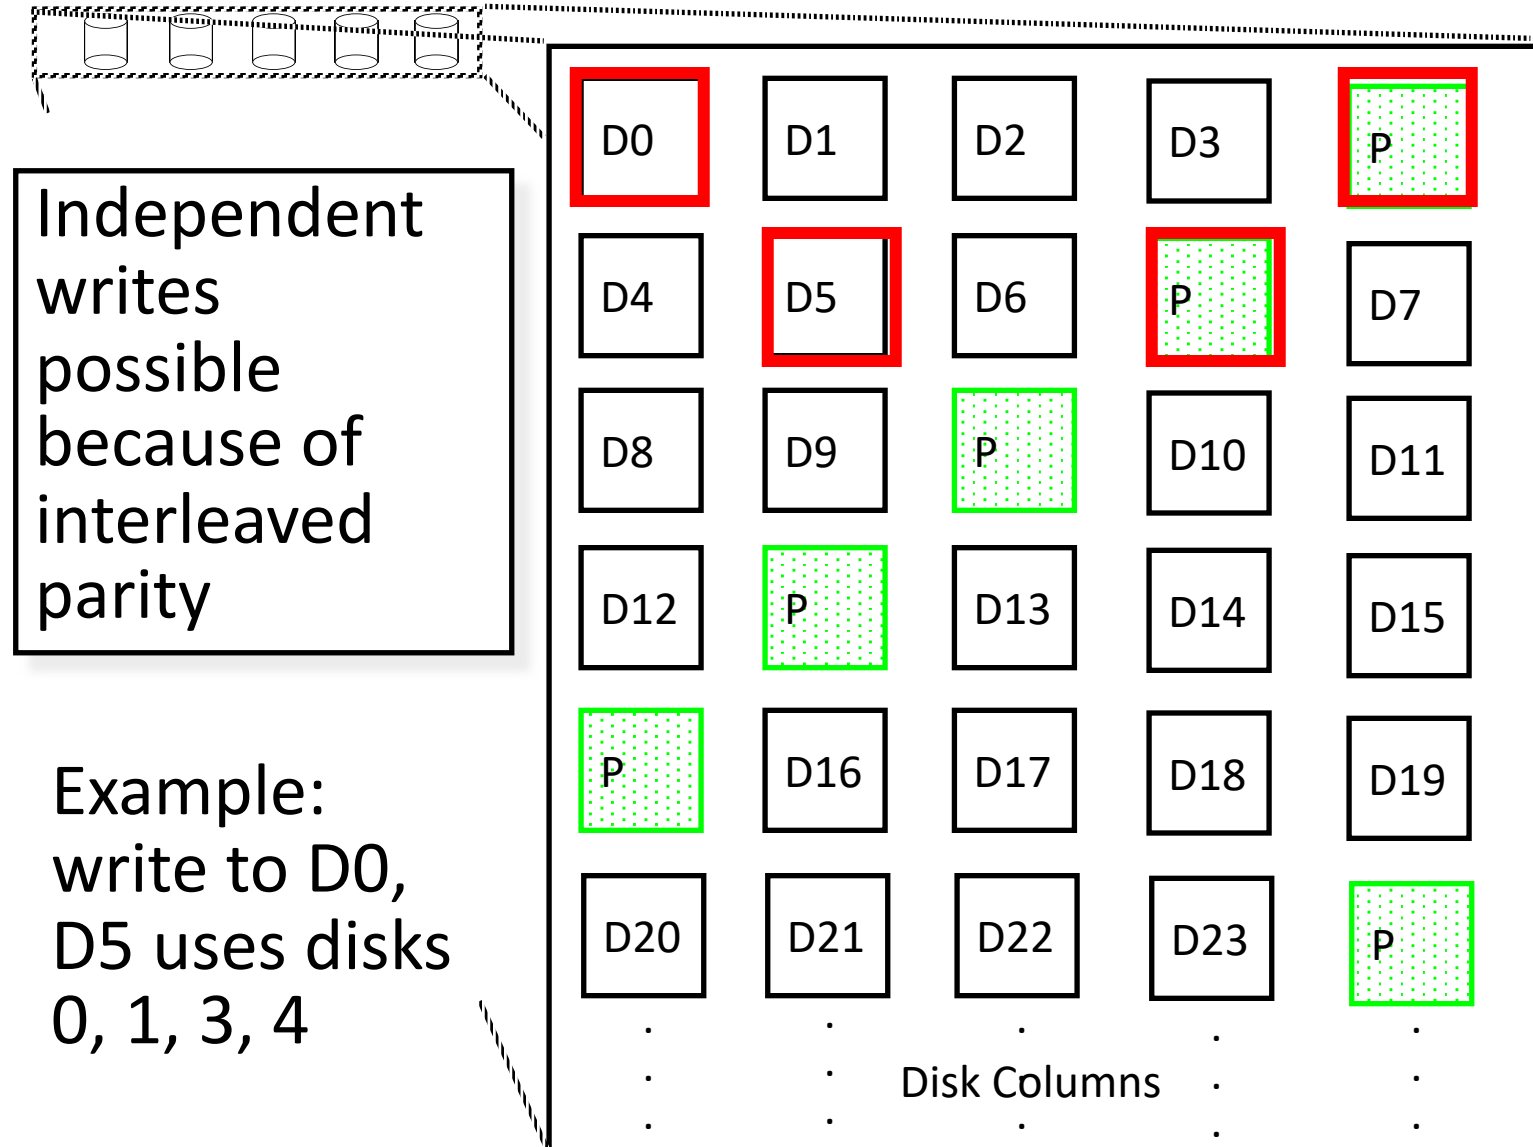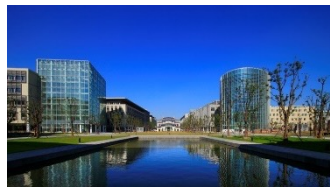Increasing Logical Disk Address

Stripe

RAID 4: High I/O Rate Parity

50

Inspiration for RAID 5

- RAID 4 works well for small reads

- Small writes (write to one disk):
  - Option 1: read other data disks, create new sum and write to Parity Disk
  - Option 2: since P has old sum, compare old data to new data, add the difference to P

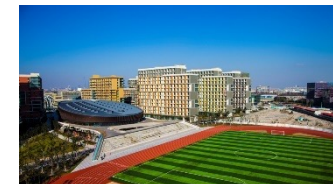- Small writes are limited by Parity Disk: Write to D0, D5 both also write to P disk
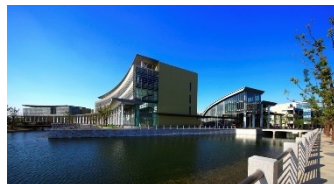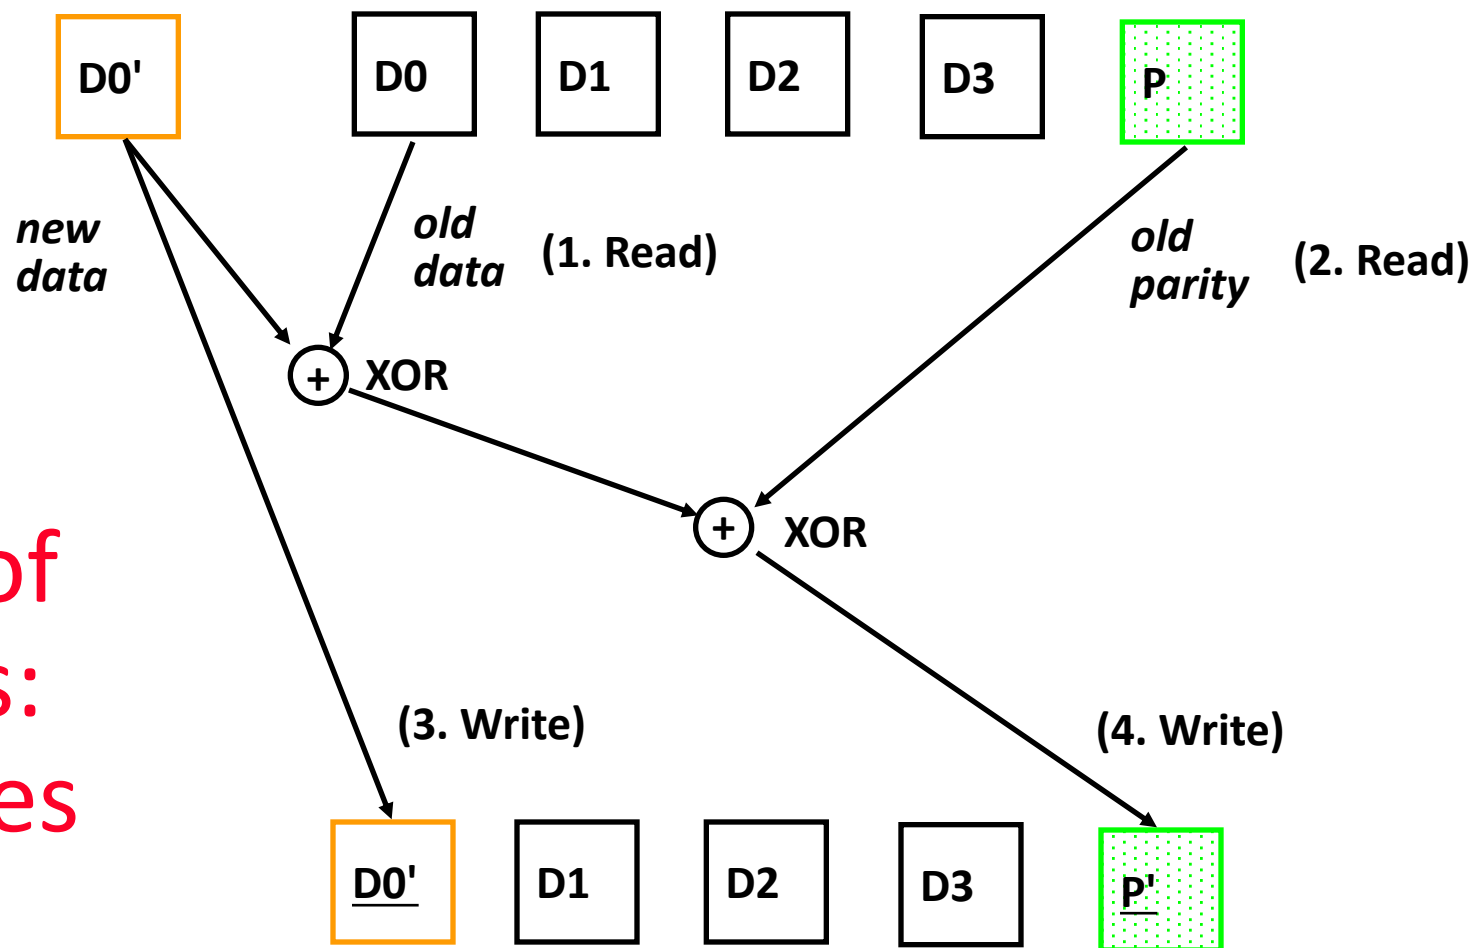
Independent writes possible because of interleaved parity

Example: write to D0, D5 uses disks 0, 1, 3, 4

Increasing Logical Disk Addresses

RAID 5: High I/O Rate Interleaved Parity

Disk Columns

52

RAID-5: *Small Write* Algorithm

1 Logical Write = 2 Physical Reads + 2 Physical Writes

D0'   D0   D1   D2   D3   P

new data   old data   (1. Read)   old parity   (2. Read)

+ XOR

+ XOR

(3. Write)   (4. Write)

D0'   D1   D2   D3   P'

Problems of Disk Arrays: Small Writes

53

And, in Conclusion, …

- Great Idea: Redundancy to Get Dependability
  - Spatial (extra hardware) and Temporal (retry if error)
- Reliability: MTTF & Annualized Failure Rate (AFR)
- Availability: % uptime
- Memory
  - Hamming ECC: correct single, detect double
- RAID
  - Interleaved data and parity